

✓ Project Title: *Home Credit Default Risk Classification using Machine Learning Techniques*

Team Information:

Group members:

1. Shoukath Ali (shshaik@iu.edu)
2. Bhargavi Vasudev Jahagirdar (bjahagir@iu.edu)
3. Palavi Dhanaji Patil (palpatil2iu.edu)
4. Saransh Kamlesh Singh (singsara@iu.edu)

Team Photo:



Phase Leadership Plan

Project Phase	Phase Description	Phase leader
Phase 0	Team creation in Canvas and Pick Project	Bhargavi Vasudev Jahagirdar
Phase 1	Project Proposal	Bhargavi Vasudev Jahagirdar
Phase 2	EDA and baseline pipeline	Palavi Dhanaji Patil
Phase 3	Final Project HCDR - feature engineering + hyperparameter tuning	Saransh Singh
Phase 4	Final Submission: Final Project HCDR	Shoukath Ali

Credit Assignment Plan

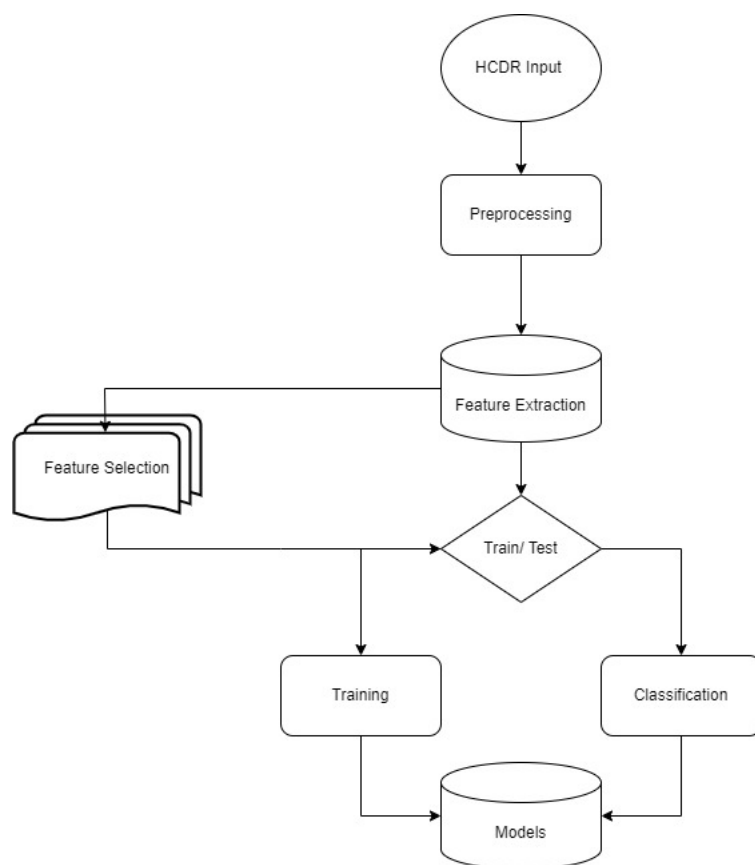
Project Phase	Phase Description	Phase Task	Owner
Phase 0	Team creation in Canvas and Pick Project	Team Creation	All
		Team Introduction	All
		Topic Selection	All
		Skillsset matrix discussion	All
		Dataset walkthrough	All
		Researching the models for the given dataset	All
		Project meetings	All
Phase 1	Project Proposal	Proposal Documentation	All
		Gantt chart	Bhargavi Jahagirdar
		Phase Leader Plan	Bhargavi Jahagirdar
		Block diagram	Bhargavi Jahagirdar
		Support Vector Machine Research	Shoukath Ali
		Random Forest Research	Saransh Singh, Palavi Patil
		Logistic Regression Research	Palavi Patil, Saransh Singh
Phase 2	EDA and baseline pipeline	Neural Networks Research	Shoukath Ali, Bhargavi Jahagirdar
		EDA - application_train.csv, bureau.csv, bureau_balance.csv, pos_cash_balance.csv, credit_cash.csv, installment_payment.csv, previous_application.csv	Bhargavi Jahagirdar, Palavi Patil, Saransh Singh
		Feature Selection	All
		Preprocessing	All
		Pipelines	Shoukath Ali
		Presentation	Saransh Singh
		Project report	Bhargavi Jahagirdar, Palavi Patil
Phase 3	Feature engineering + Hyperparameter tuning	Kaggle Submission	Palavi Patil
		Feature Engineering	Saransh Singh
		Hyperparameter tuning	Saransh Singh
		Additional Feature Engineering	Palavi Patil, Bhargavi Jahagirdar
		Model training: Random Forest, SVM, Logistic Regression	Saransh Singh, Palavi Patil, Bhargavi Jahagirdar
		Deploying the final model: Random Forest, SVM, Logistic Regression	Saransh Singh, Palavi Patil, Bhargavi Jahagirdar
		Initial Neural Networks Implementation	Shoukath Ali
Phase 4	Final Project HCDR / CaDoD - PyTorch Deep Learning	Project report	Palavi Patil, Bhargavi Jahagirdar
		Presentation	Saransh Singh
		Improving the existing models	All
		Implementation of Multi Layer Perceptron model (MLP)	All
		Kaggle Submission	Shoukath Ali
		Project Report	All
		Presentation	All

✓ Abstract

The given problem is the Home Credit Default Risk Prediction presented by Kaggle. The team aims to provide suitable Machine Learning models that can accurately predict the results for the given problem statement. For this binary classification problem, we used a Random Forest Classifier, Logistic Regression, and Neural Networks. We performed the project in 4 phases. The previous phases involved creating the project proposal, performing EDA, creating baseline pipelines, feature engineering, and hyperparameter tuning. The current phase involves the implementation of Neural Networks and using Neural Network architectures alongside the implementation of Random Forest Classifier, and Logistic Regression. Accuracy for each of them is 91.95% and 91.7%, respectively. We obtained the best results through the Random Forest model. We submitted the below models to the Kaggle competition and achieved a public score of 73.409% using Neural Networks. We used two architectures for Neural Networks and achieved the accuracy of 91 % and 91 % for each of them

✓ Data Description

Data File Name	Description
application_{train test}.csv	Main table for training, includes static data for loan applications. Each row represents a loan in the sample.
bureau.csv	Contains information about clients' previous credits reported to Credit Bureau for those with a loan in the sample. Multiple rows for each loan corresponding to the client's previous credits.
bureau_balance.csv	Monthly balances of previous credits in Credit Bureau, with one row per month for each previous credit, resulting in many rows.
POS_CASH_balance.csv	Monthly balance snapshots of previous POS and cash loans with Home Credit, featuring one row per month for each previous credit related to loans in the sample.
credit_card_balance.csv	Monthly balance snapshots of previous credit cards with Home Credit, with one row per month for each previous credit card related to loans in the sample.
previous_application.csv	Contains information about all previous loan applications for Home Credit clients in the sample, with one row for each previous application.
installments_payments.csv	Provides repayment history for previously disbursed credits from Home Credit related to the loans in the sample. Each row corresponds to a payment made or a missed payment.
HomeCredit_columns_description.csv	Contains column descriptions for the various data files.



Machine Learning Models and Pipelines for the Dataset

4.1 SVM Metrics used 1) Classification metrics (F1 score, recall, precision) 2) accuracy_score

The problem statement's output is to predict the binary outcomes (classification problem). SVM is capable of improving the accuracy of the ML model by considering the proper loss function(Margin) or best parameters. Required Pipelines Standardize, Normalize, Impute (in case of missing data), stratified K fold (for skewed data), feature engineering, Grid Search CV (hyperparameter tuning), SVM, and Model evaluation. We will build the final pipelines using multiple sub-pipelines. For Example, numerical_pipeline = Pipeline([('imputer', SimpleImputer(strategy='mean')), ('scaler', StandardScaler()), ('pca', PCA(n_components=0.95))])

categorical_pipeline = Pipeline([('imputer', SimpleImputer(strategy='most_frequent')), ('one_hot', OneHotEncoder())])

preprocessor = ColumnTransformer([('num', numerical_pipeline, numerical_cols), ('cat', categorical_pipeline, categorical_cols)])

4.2 Random Forest

For the Home Credit Default Risk dataset, where the goal is typically to predict the loan repayment abilities of borrowers, a Random Forest model can be an excellent choice due to its versatility, robustness to overfitting, and ability to handle imbalanced datasets. Here's a proposal outline that answers your questions within a tight character limit:

Metrics used 1) AUC-ROC: Ideal for classification on imbalanced datasets, measures the ability to distinguish between classes.

2) F1-Score: Balances precision and recall, useful when false negatives and false positives are crucial.

3) Accuracy: Provides a quick understanding of overall performance, though less informative on its own for imbalanced data. Why is your model best for the problem statement Random Forest is robust to overfitting and excellent for handling complex datasets with many features, like Home Credit. It's also good for imbalanced classes and does not require scaling of data.

Required Pipelines 1) Data Preprocessing Pipeline: To handle missing values, encode categorical variables, and potentially scale features if necessary.

2) Model Training Pipeline: For fitting the Random Forest model with cross-validation to avoid overfitting.

3) Evaluation Pipeline: To apply the chosen metrics and validate the model's performance. Other pipelines:

4) Feature Engineering Pipeline: To create new features that can provide additional insights for the Random Forest algorithm.

5) Model Optimization Pipeline: Using grid search or random search to fine-tune hyperparameters for the Random Forest. For Example, `feature_engineering_pipeline = Pipeline([('polynomial_features', PolynomialFeatures(degree=2, include_bias=False)), ('feature_selection', SelectFromModel(RandomForestClassifier(n_estimators=100)))])`

4.3 Logistic Regression Metrics used 1) Confusion Matrix: The confusion matrix will provide a detail of the model's predictions, which will include true positives, true negatives, false positives, and false negatives. 2) Accuracy Score: Accuracy will measure the proportion of correct predictions made by the logistic regression model. 3) Precision Score: Precision will measure the accuracy of positive predictions made by the model. It will tell us the proportion of positive predictions that were correct. 4) F1 score: The F1 score will use precision and recall. It balances the trade-off between precision and recall. Why is your model best for the problem statement For the HCDR dataset, we have to classify whether the client can repay the loan money. The classes will be Yes or No. For this, we can use logistic regression as it is specifically designed for binary classification problems. Logistic regression is a relatively simple model, which is computationally efficient and can be trained quickly.

Required Pipelines 1) Standardization: This preprocessing step is required to do the scaling of the input features. 2) Imputation: Multiple methods can be used to impute missing data in our dataset. We can replace missing values with the mean. 3) GridSearchCV: This will be used to tune the hyperparameters so we can find the best combination that will optimize our performance. 4) Model pipeline: Here we can combine our model with other pipelines. For Example, `preprocessingPipeline = Pipeline([('imputer', SimpleImputer(strategy='mean')), ('scaler', StandardScaler()),])` Above metrics and pipelines can be used to train HCDR dataset on Logistic Regression Model.

4.4 Neural Networks Metrics used 1) Confusion matrix 2) Classification metrics (F1 score, recall, precision) 3) accuracy_score Why is your model best for the problem statement 1) Scalability: Neural networks can be scaled up with more layers and neurons to handle increasingly complex classification tasks. Deep neural networks (deep learning) are particularly effective at capturing hierarchical representations in the data, which can improve classification accuracy. 2) Regularization Techniques: Neural networks offer various regularization techniques to prevent overfitting, including dropout, weight decay, and early stopping, which can help improve generalization on classification tasks. 3) Neural networks can automatically learn relevant features from raw data. Instead of hand-crafting features, as is often done in traditional machine learning, neural networks learn representations from the data during training.

Required Pipelines 1) Standardize 2) Normalize, Impute (in case of missing data) 3) stratified K fold (for skewed data), data augmentation/pre-processing 4) Neural Network Model 5) Grid Search CV (hyperparameter tuning), Optimizing/Regularisation (activation function, loss function, epoch, and batch values) 6) Model evaluation.

Other pipelines: we build the final pipeline using all the sub-pipelines, and here, feature engineering is not required as Neural networks are capable of doing feature learning from the data itself.

F1 Score: The F1 score is a metric used in binary classification to measure a balance between precision and recall. It's calculated using the following formula:

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

AUC-ROC (Area Under the Receiver Operating Characteristic Curve): The AUC-ROC is a performance measurement for classification problems at various threshold settings. The ROC curve is a plot of the true positive rate (TPR) against the false positive rate (FPR) at various threshold values. The AUC-ROC represents the area under this ROC curve.

$$TPR = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

$$FPR = \frac{\text{False Positives}}{\text{False Positives} + \text{True Negatives}}$$

✓ Project Tasks

Project Phase	Phase Description	Project Tasks	Week 11	Week 12	Week 13	Week 14	Week 15	Week 16
Phase 0	Team creation in Canvas and Pick Project	Team Creation						
		Project Topic Selection						
Phase 1	Project Proposal	Proposal Documentation : Project description, data, metrics, baseline models, baseline pipeline, and other planned pipelines						
		Project Plan creation						
		Phase leader plan creation						
		Credit Assignment Plan						
Phase 2	EDA and baseline pipeline	Exploratory Data Analysis						
		Baseline model creation						
		Feature Engineering						
		Hyperparameter tuning						
		Phase documentation						
Phase 3	Final Project HCDR - feature engineering + hyperparameter tuning	Presentation video preparation						
		Feature Engineering						
		Hyperparameter tuning						
		Additional Feature Engineering						
		Model training						
Phase 4	Final Submission: Final Project HCDR	Deploying the final model						
		Neural Networks Implementation						
		Model training						
		Deploying the final model						
		Final project report						
		Presentation video preparation						

✓ Importing Datasets

```
! pip install -q kaggle
```

```
from google.colab import files
```

```
files.upload()
```

Choose Files kaggle.json

- **kaggle.json**(application/json) - 72 bytes, last modified: 11/14/2023 - 100% done

Saving kaggle.json to kaggle.json

```
{'kaggle.json': b'{"username": "saranshsingh2626", "key": "5c9f084b0d4b8f05324de44a8f9c3b97"}'}
```

```
! mkdir ~/.kaggle
```

```
! cp kaggle.json ~/.kaggle/
```

```
! chmod 600 ~/.kaggle/kaggle.json
```

```
! kaggle datasets list
```

ref	title	size	lastUpdate
thedrcat/daigt-v2-train-dataset	DAIGT V2 Train Dataset	29MB	2023-11-
muhammadbinimran/housing-price-prediction-data	Housing Price Prediction Data	763KB	2023-11-
carlmcbriedeellis/llm-7-prompt-training-dataset	LLM: 7 prompt training dataset	41MB	2023-11-
thedrcat/daigt-proper-train-dataset	DAIGT Proper Train Dataset	119MB	2023-11-
joebeachcapital/30000-spotify-songs	30000 Spotify Songs	3MB	2023-11-
jacksondivakarr/laptop-price-prediction-dataset	Laptop Price Prediction Dataset	119KB	2023-11-
ddosad/auto-sales-data	Automobile Sales data	79KB	2023-11-
julnazz/diabetes-health-indicators-dataset	Diabetes Health Indicators Dataset	5MB	2023-11-
stealthtechnologies/predict-lifespan-of-a-comet-goldfish	Predict lifespan of a comet goldfish	25KB	2023-11-
nelgiriyeewithana/world-educational-data	World Educational Data	9KB	2023-11-
thedevasator/bank-term-deposit-predictions	Bank Term Deposit Predictions	541KB	2023-11-
maso0dahmed/video-games-data	Video Games Data	5MB	2023-11-
alejopaullier/daigt-external-dataset	DAIGT External Dataset	3MB	2023-10-

nelgiriyeewithana/australian-vehicle-prices	Australian Vehicle Prices	582KB	2023-11-
prasad22/healthcare-dataset	Healthcare Dataset ✓	483KB	2023-
adampq/linkedin-jobs-machine-learning-data-set	LinkedIn Job Postings - Machine Learning Data Set	38MB	2023-11-
jacksondivakarr/online-shopping-dataset	Online Shopping Dataset 🇮🇹🇺🇸🇬🇧	5MB	202
asimislam/30-yrs-stock-market-data	30 yrs Stock Market Data	882KB	2023-11-
imtkaggleteam/life-expectancy	Life Expectancy	730KB	2023-11-
sujoykapadnis/products-datasets	Detailed Products Datasets	100KB	2023-11-

#Downloading the dataset from Kaggle

! kaggle competitions download -c home-credit-default-risk

Downloading home-credit-default-risk.zip to /content

98% 673M/688M [00:08<00:00, 137MB/s]

100% 688M/688M [00:08<00:00, 88.4MB/s]

#Creating a directory for the adding the dataset

! mkdir dataset

unzip the downloaded dataset file in the dataset directory

! unzip home-credit-default-risk.zip -d dataset

Archive: home-credit-default-risk.zip

inflating: dataset/HomeCredit_columns_description.csv

inflating: dataset/POS_CASH_balance.csv

inflating: dataset/application_test.csv

inflating: dataset/application_train.csv

inflating: dataset/bureau.csv

inflating: dataset/bureau_balance.csv

inflating: dataset/credit_card_balance.csv

inflating: dataset/installments_payments.csv

inflating: dataset/previous_application.csv

inflating: dataset/sample_submission.csv

#Loading the data files in dataframes

import numpy as np

import pandas as pd

df_application_train = pd.read_csv('dataset/application_train.csv')

df_application_test = pd.read_csv('dataset/application_test.csv')

df_bureau = pd.read_csv('dataset/bureau.csv')

df_bureau_balance = pd.read_csv('dataset/bureau_balance.csv')

df_pos_cash_balance = pd.read_csv('dataset/POS_CASH_balance.csv')

df_credit_card_balance = pd.read_csv('dataset/credit_card_balance.csv')

df_previous_application = pd.read_csv('dataset/previous_application.csv')

df_installments_payments = pd.read_csv('dataset/installments_payments.csv')

✓ Data Description

✓ Data File: application_train.csv

File description:

- This is the main table, broken into two files for Train (with TARGET) and Test (without TARGET).
- Static data for all applications. One row represents one loan in our data sample.

```
print("File name: application_train.csv")
```

```
print("Number of rows, columns:", df_application_train.shape)
```

```
print("Number of Missing Values: " + str(df_application_train.isna().sum().sum()))
```

```
df_application_train.head(5)
```

File name: application_train.csv
 Number of rows, columns: (307511, 122)
 Number of Missing Values: 9152465

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CRE
0	100002	1	Cash loans	M	N	Y	0	202500.0	4065
1	100003	0	Cash loans	F	N	N	0	270000.0	12935
2	100004	0	Revolving loans	M	Y	Y	0	67500.0	1350

✓ Data File: application_test.csv

File description:

- This is the main table, broken into two files for Train (with TARGET) and Test (without TARGET).
- Static data for all applications. One row represents one loan in our data sample.

```
print("File name: application_test.csv")
print("Number of rows, columns:", df_application_test.shape)
print("Number of Missing Values: " + str(df_application_test.isna().sum().sum()))
df_application_test.head(5)
```

File name: application_test.csv
 Number of rows, columns: (48744, 121)
 Number of Missing Values: 1404419

	SK_ID_CURR	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT
0	100001	Cash loans	F	N	Y	0	135000.0	568800.0	
1	100005	Cash loans	M	N	Y	0	99000.0	222768.0	
2	100013	Cash loans	M	Y	Y	0	202500.0	663264.0	
3	100028	Cash loans	F	N	Y	2	315000.0	1575000.0	
4	100038	Cash loans	M	Y	N	1	180000.0	625500.0	

5 rows x 121 columns

✓ Data File: bureau.csv

File description:

- All client's previous credits provided by other financial institutions that were reported to Credit Bureau (for clients who have a loan in our sample).
- For every loan in our sample, there are as many rows as number of credits the client had in Credit Bureau before the application date.

```
print("File name: bureau.csv")
print("Number of rows, columns:", df_bureau.shape)
print("Number of Missing Values: " + str(df_bureau.isna().sum().sum()))
df_bureau.head(5)
```

File name: bureau.csv
 Number of rows, columns: (1716428, 17)
 Number of Missing Values: 3939947

	SK_ID_CURR	SK_ID_BUREAU	CREDIT_ACTIVE	CREDIT_CURRENCY	DAYS_CREDIT	CREDIT_DAY_OVERDUE	DAYS_CREDIT_ENDDATE	DAYS_ENDDA1
0	215354	5714462	Closed	currency 1	-497	0	-153.0	
1	215354	5714463	Active	currency 1	-208	0	1075.0	
2	215354	5714464	Active	currency 1	-203	0	528.0	
3	215354	5714465	Active	currency 1	-203	0	NaN	
4	215354	5714466	Active	currency 1	-629	0	1197.0	

✓ Data File: bureau_balance.csv

File description:

- Monthly balances of previous credits in Credit Bureau.
- This table has one row for each month of history of every previous credit reported to Credit Bureau – i.e the table has (#loans in sample * # of relative previous credits * # of months where we have some history observable for the previous credits) rows.

```
print("File name: bureau_balance.csv")
print("Number of rows, columns:", df_bureau_balance.shape)
print("Number of Missing Values: " + str(df_bureau_balance.isna().sum().sum()))
df_bureau_balance.head(5)
```

```
File name: bureau_balance.csv
Number of rows, columns: (27299925, 3)
Number of Missing Values: 0
```

	SK_ID_BUREAU	MONTHS_BALANCE	STATUS
0	5715448	0	C
1	5715448	-1	C
2	5715448	-2	C
3	5715448	-3	C
4	5715448	-4	C

✓ Data File: pos_cash_balance.csv

File description:

- Monthly balance snapshots of previous POS (point of sales) and cash loans that the applicant had with Home Credit.
- This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credits * # of months in which we have some history observable for the previous credits) rows.

```
print("File name: pos_cash_balance.csv")
print("Number of rows, columns:", df_pos_cash_balance.shape)
print("Number of Missing Values: " + str(df_pos_cash_balance.isna().sum().sum()))
df_pos_cash_balance.head(5)
```

```
File name: pos_cash_balance.csv
Number of rows, columns: (10001358, 8)
Number of Missing Values: 52158
```

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	CNT_INSTALMENT	CNT_INSTALMENT_FUTURE	NAME_CONTRACT_STATUS	SK_DPD	SK_DPD_DEF
0	1803195	182943	-31	48.0	45.0	Active	0	0
1	1715348	367990	-33	36.0	35.0	Active	0	0
2	1784872	397406	-32	12.0	9.0	Active	0	0
3	1903291	269225	-35	48.0	42.0	Active	0	0
4	2341044	334279	-35	36.0	35.0	Active	0	0

✓ Data File: credit_card_balance.csv

File description:

- Monthly balance snapshots of previous credit cards that the applicant has with Home Credit.
- This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credit cards * # of months where we have some history observable for the previous credit card) rows.

```
print("File name: credit_card_balance.csv")
print("Number of rows, columns:", df_credit_card_balance.shape)
print("Number of Missing Values: " + str(df_credit_card_balance.isna().sum().sum()))
df_credit_card_balance.head(5)
```

File name: credit_card_balance.csv
 Number of rows, columns: (3840312, 23)
 Number of Missing Values: 5877356

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	AMT_BALANCE	AMT_CREDIT_LIMIT_ACTUAL	AMT_DRAWINGS_ATM_CURRENT	AMT_DRAWINGS_CURREN
0	2562384	378907	-6	56.970	135000	0.0	877
1	2582071	363914	-1	63975.555	45000	2250.0	2250
2	1740877	371185	-7	31815.225	450000	0.0	(
3	1389973	337855	-4	236572.110	225000	2250.0	2250
4	1891521	126868	-1	453919.455	450000	0.0	11547

5 rows × 23 columns

✓ Data File: previous_application.csv

File description:

- All previous applications for Home Credit loans of clients who have loans in our sample.
- There is one row for each previous application related to loans in our data sample.

```
print("File name: df_previous_application.csv")
print("Number of rows, columns:", df_previous_application.shape)
print("Number of Missing Values: " + str(df_previous_application.isna().sum().sum()))
df_previous_application.head(5)
```

File name: df_previous_application.csv
 Number of rows, columns: (1670214, 37)
 Number of Missing Values: 11109336

	SK_ID_PREV	SK_ID_CURR	NAME_CONTRACT_TYPE	AMT_ANNUITY	AMT_APPLICATION	AMT_CREDIT	AMT_DOWN_PAYMENT	AMT_GOODS_PRICE	WE
0	2030495	271877	Consumer loans	1730.430	17145.0	17145.0	0.0	17145.0	
1	2802425	108129	Cash loans	25188.615	607500.0	679671.0	NaN	607500.0	
2	2523466	122040	Cash loans	15060.735	112500.0	136444.5	NaN	112500.0	
3	2819243	176158	Cash loans	47041.335	450000.0	470790.0	NaN	450000.0	
4	1784265	202054	Cash loans	31924.395	337500.0	404055.0	NaN	337500.0	

5 rows × 37 columns

✓ Data File: installments_payments.csv

File description:

- Repayment history for the previously disbursed credits in Home Credit related to the loans in our sample.
- There is a) one row for every payment that was made plus b) one row each for missed payment.
- One row is equivalent to one payment of one installment OR one installment corresponding to one payment of one previous Home Credit credit related to loans in our sample.

```
print("File name: installments_payments.csv")
print("Number of rows, columns:", df_installments_payments.shape)
print("Number of Missing Values: " + str(df_installments_payments.isna().sum().sum()))
df_installments_payments.head(5)
```

File name: installments_payments.csv

Methods

-	-	-	-	-	-	-	-	-	-
---	---	---	---	---	---	---	---	---	---

Discussion:

1	1330831	151639	0.0	34	-2156.0	-2156.0	1716.525
---	---------	--------	-----	----	---------	---------	----------

Exploratory Data Analysis

Data File: application_test.csv

File description:

- This is the main table, broken into two files for Train (with TARGET) and Test (without TARGET).
- Static data for all applications. One row represents one loan in our data sample.

```
print("File name: application_test.csv")
print("Number of rows, columns:", df_application_test.shape)
print("Number of Missing Values: " + str(df_application_test.isna().sum().sum()))
df_application_test.head(5)
```

File name: application_test.csv
Number of rows, columns: (48744, 121)
Number of Missing Values: 1404419

	SK_ID_CURR	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT
0	100001	Cash loans	F	N	Y	0	135000.0	568800.0	
1	100005	Cash loans	M	N	Y	0	99000.0	222768.0	
2	100013	Cash loans	M	Y	Y	0	202500.0	663264.0	
3	100028	Cash loans	F	N	Y	2	315000.0	1575000.0	
4	100038	Cash loans	M	Y	N	1	180000.0	625500.0	

5 rows x 121 columns

Data File: bureau.csv

File description:

- All client's previous credits provided by other financial institutions that were reported to Credit Bureau (for clients who have a loan in our sample).
- For every loan in our sample, there are as many rows as number of credits the client had in Credit Bureau before the application date.

```
print("File name: bureau.csv")
print("Number of rows, columns:", df_bureau.shape)
print("Number of Missing Values: " + str(df_bureau.isna().sum().sum()))
df_bureau.head(5)
```

File name: bureau.csv
Number of rows, columns: (1716428, 17)
Number of Missing Values: 3939947

	SK_ID_CURR	SK_ID_BUREAU	CREDIT_ACTIVE	CREDIT_CURRENCY	DAYS_CREDIT	CREDIT_DAY_OVERDUE	DAYS_CREDIT_ENDDATE	DAYS_ENDDATE
0	215354	5714462	Closed	currency 1	-497	0	-153.0	
1	215354	5714463	Active	currency 1	-208	0	1075.0	
2	215354	5714464	Active	currency 1	-203	0	528.0	
3	215354	5714465	Active	currency 1	-203	0	NaN	
4	215354	5714466	Active	currency 1	-629	0	1197.0	

✓ Data File: bureau_balance.csv

File description:

- Monthly balances of previous credits in Credit Bureau.
- This table has one row for each month of history of every previous credit reported to Credit Bureau – i.e the table has (#loans in sample * # of relative previous credits * # of months where we have some history observable for the previous credits) rows.

```
print("File name: bureau_balance.csv")
print("Number of rows, columns:", df_bureau_balance.shape)
print("Number of Missing Values: " + str(df_bureau_balance.isna().sum().sum()))
df_bureau_balance.head(5)
```

```
File name: bureau_balance.csv
Number of rows, columns: (27299925, 3)
Number of Missing Values: 0
```

	SK_ID_BUREAU	MONTHS_BALANCE	STATUS
0	5715448	0	C
1	5715448	-1	C
2	5715448	-2	C
3	5715448	-3	C
4	5715448	-4	C

✓ Data File: pos_cash_balance.csv

File description:

- Monthly balance snapshots of previous POS (point of sales) and cash loans that the applicant had with Home Credit.
- This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credits * # of months in which we have some history observable for the previous credits) rows.

```
print("File name: pos_cash_balance.csv")
print("Number of rows, columns:", df_pos_cash_balance.shape)
print("Number of Missing Values: " + str(df_pos_cash_balance.isna().sum().sum()))
df_pos_cash_balance.head(5)
```

```
File name: pos_cash_balance.csv
Number of rows, columns: (10001358, 8)
Number of Missing Values: 52158
```

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	CNT_INSTALMENT	CNT_INSTALMENT_FUTURE	NAME_CONTRACT_STATUS	SK_DPD	SK_DPD_DEF
0	1803195	182943	-31	48.0	45.0	Active	0	0
1	1715348	367990	-33	36.0	35.0	Active	0	0
2	1784872	397406	-32	12.0	9.0	Active	0	0
3	1903291	269225	-35	48.0	42.0	Active	0	0
4	2341044	334279	-35	36.0	35.0	Active	0	0

✓ Data File: credit_card_balance.csv

File description:

- Monthly balance snapshots of previous credit cards that the applicant has with Home Credit.
- This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credit cards * # of months where we have some history observable for the previous credit card) rows.

```
print("File name: credit_card_balance.csv")
print("Number of rows, columns:", df_credit_card_balance.shape)
print("Number of Missing Values: " + str(df_credit_card_balance.isna().sum().sum()))
df_credit_card_balance.head(5)
```

File name: credit_card_balance.csv
 Number of rows, columns: (3840312, 23)
 Number of Missing Values: 5877356

	SK_ID_PREV	SK_ID_CURR	MONTHS_BALANCE	AMT_BALANCE	AMT_CREDIT_LIMIT_ACTUAL	AMT_DRAWINGS_ATM_CURRENT	AMT_DRAWINGS_CURREN
0	2562384	378907	-6	56.970	135000	0.0	877
1	2582071	363914	-1	63975.555	45000	2250.0	2250
2	1740877	371185	-7	31815.225	450000	0.0	(
3	1389973	337855	-4	236572.110	225000	2250.0	2250
4	1891521	126868	-1	453919.455	450000	0.0	11547

5 rows × 23 columns

✓ Data File: previous_application.csv

File description:

- All previous applications for Home Credit loans of clients who have loans in our sample.
- There is one row for each previous application related to loans in our data sample.

```
print("File name: df_previous_application.csv")
print("Number of rows, columns:", df_previous_application.shape)
print("Number of Missing Values: " + str(df_previous_application.isna().sum().sum()))
df_previous_application.head(5)
```

File name: df_previous_application.csv
 Number of rows, columns: (1670214, 37)
 Number of Missing Values: 11109336

	SK_ID_PREV	SK_ID_CURR	NAME_CONTRACT_TYPE	AMT_ANNUITY	AMT_APPLICATION	AMT_CREDIT	AMT_DOWN_PAYMENT	AMT_GOODS_PRICE	WE
0	2030495	271877	Consumer loans	1730.430	17145.0	17145.0	0.0	17145.0	
1	2802425	108129	Cash loans	25188.615	607500.0	679671.0	NaN	607500.0	
2	2523466	122040	Cash loans	15060.735	112500.0	136444.5	NaN	112500.0	
3	2819243	176158	Cash loans	47041.335	450000.0	470790.0	NaN	450000.0	
4	1784265	202054	Cash loans	31924.395	337500.0	404055.0	NaN	337500.0	

5 rows × 37 columns

✓ Data File: installments_payments.csv

File description:

- Repayment history for the previously disbursed credits in Home Credit related to the loans in our sample.
- There is a) one row for every payment that was made plus b) one row each for missed payment.
- One row is equivalent to one payment of one installment OR one installment corresponding to one payment of one previous Home Credit credit related to loans in our sample.

```
print("File name: installments_payments.csv")
print("Number of rows, columns:", df_installments_payments.shape)
print("Number of Missing Values: " + str(df_installments_payments.isna().sum().sum()))
df_installments_payments.head(5)
```

File names: installments payments csv

File: application_train.csv

```
print("Number of Rows: " + str(df_application_train.shape[0]))
print("Number of Columns: " + str(df_application_train.shape[1]))
print("Number of Total Missing Values: " + str(df_application_train.isna().sum().sum()))
print("Data Frame Shape: " + str(df_application_train.shape))
print("Number of Missing Values by Feature: " + str(df_application_train.isna().sum()))
print("Data Types:", df_application_train.dtypes)
print("Data Frame: Data Types", df_application_train.dtypes.value_counts())
print("Summary", df_application_train.describe())
print("Correlation Statistic", df_application_train.corr())
print("Summary Information:", df_application_train.info())
```

```
SK_ID_CURR      0.002099
TARGET          0.000788
CNT_CHILDREN   -0.002436
AMT_INCOME_TOTAL 0.002387
AMT_CREDIT     -0.001275
...
AMT_REQ_CREDIT_BUREAU_DAY  0.217412
AMT_REQ_CREDIT_BUREAU_WEEK 1.000000
AMT_REQ_CREDIT_BUREAU_MON -0.014096
AMT_REQ_CREDIT_BUREAU_QRT -0.015115
AMT_REQ_CREDIT_BUREAU_YEAR 0.018917
```

```
AMT_REQ_CREDIT_BUREAU_MON \
SK_ID_CURR      0.000485
TARGET         -0.012462
CNT_CHILDREN   -0.010808
AMT_INCOME_TOTAL 0.024700
AMT_CREDIT      0.054451
...
AMT_REQ_CREDIT_BUREAU_DAY -0.005258
AMT_REQ_CREDIT_BUREAU_WEEK -0.014096
AMT_REQ_CREDIT_BUREAU_MON 1.000000
AMT_REQ_CREDIT_BUREAU_QRT -0.007789
AMT_REQ_CREDIT_BUREAU_YEAR -0.004975
```

```
AMT_REQ_CREDIT_BUREAU_QRT \
SK_ID_CURR      0.001025
TARGET         -0.002022
CNT_CHILDREN   -0.007836
AMT_INCOME_TOTAL 0.004859
AMT_CREDIT      0.015925
...
AMT_REQ_CREDIT_BUREAU_DAY -0.004416
AMT_REQ_CREDIT_BUREAU_WEEK -0.015115
AMT_REQ_CREDIT_BUREAU_MON -0.007789
AMT_REQ_CREDIT_BUREAU_QRT 1.000000
AMT_REQ_CREDIT_BUREAU_YEAR 0.076208
```

```
AMT_REQ_CREDIT_BUREAU_YEAR
SK_ID_CURR      0.004659
TARGET          0.019930
CNT_CHILDREN   -0.041550
AMT_INCOME_TOTAL 0.011690
AMT_CREDIT     -0.048448
...
AMT_REQ_CREDIT_BUREAU_DAY -0.003355
AMT_REQ_CREDIT_BUREAU_WEEK 0.018917
AMT_REQ_CREDIT_BUREAU_MON -0.004975
AMT_REQ_CREDIT_BUREAU_QRT 0.076208
AMT_REQ_CREDIT_BUREAU_YEAR 1.000000
```

```
[106 rows x 106 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 286.2+ MB
Summary Information: None
```

File: application_test.csv

```

print("Number of Rows: " + str(df_application_test.shape[0]))
print("Number of Columns: " + str(df_application_test.shape[1]))
print("Number of Total Missing Values: " + str(df_application_test.isna().sum().sum()))
print("Data Frame Shape: " + str(df_application_test.shape))
print("Number of Missing Values by Feature: " + str(df_application_test.isna().sum()))
print("Data Types:",df_application_test.dtypes)
print("Data Frame: Data Types", df_application_test.dtypes.value_counts())
print("Summary", df_application_test.describe())
print("Correlation Statistic", df_application_test.corr())
print("Summary Information:", df_application_test.info())

```

```

SK_ID_CURR      0.001178
CNT_CHILDREN    0.007523
AMT_INCOME_TOTAL -0.002867
AMT_CREDIT      0.002904
AMT_ANNUITY     0.003085
...
AMT_REQ_CREDIT_BUREAU_DAY    0.035567
AMT_REQ_CREDIT_BUREAU_WEEK  1.000000
AMT_REQ_CREDIT_BUREAU_MON   0.054291
AMT_REQ_CREDIT_BUREAU_QRT   0.024957
AMT_REQ_CREDIT_BUREAU_YEAR  -0.000252

```

```

                                AMT_REQ_CREDIT_BUREAU_MON \
SK_ID_CURR      0.000430
CNT_CHILDREN    -0.008337
AMT_INCOME_TOTAL 0.008691
AMT_CREDIT      -0.000156
AMT_ANNUITY     0.005695
...
AMT_REQ_CREDIT_BUREAU_DAY    0.005877
AMT_REQ_CREDIT_BUREAU_WEEK  0.054291
AMT_REQ_CREDIT_BUREAU_MON   1.000000
AMT_REQ_CREDIT_BUREAU_QRT   0.005446
AMT_REQ_CREDIT_BUREAU_YEAR  0.026118

```

```

                                AMT_REQ_CREDIT_BUREAU_QRT \
SK_ID_CURR      -0.002092
CNT_CHILDREN     0.029006
AMT_INCOME_TOTAL 0.007410
AMT_CREDIT      -0.007750
AMT_ANNUITY     0.012443
...
AMT_REQ_CREDIT_BUREAU_DAY    0.006509
AMT_REQ_CREDIT_BUREAU_WEEK  0.024957
AMT_REQ_CREDIT_BUREAU_MON   0.005446
AMT_REQ_CREDIT_BUREAU_QRT   1.000000
AMT_REQ_CREDIT_BUREAU_YEAR  -0.013081

```

```

                                AMT_REQ_CREDIT_BUREAU_YEAR
SK_ID_CURR      0.003457
CNT_CHILDREN    -0.039265
AMT_INCOME_TOTAL 0.003281
AMT_CREDIT      -0.034533
AMT_ANNUITY     -0.044901
...
AMT_REQ_CREDIT_BUREAU_DAY    0.002002
AMT_REQ_CREDIT_BUREAU_WEEK  -0.000252
AMT_REQ_CREDIT_BUREAU_MON   0.026118
AMT_REQ_CREDIT_BUREAU_QRT  -0.013081
AMT_REQ_CREDIT_BUREAU_YEAR  1.000000

```

```

[105 rows x 105 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48744 entries, 0 to 48743
Columns: 121 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(40), object(16)
memory usage: 45.0+ MB
Summary Information: None

```

✓ **File:** bureau.csv

```

print("Number of Rows: " + str(df_bureau.shape[0]))
print("Number of Columns: " + str(df_bureau.shape[1]))
print("Number of Total Missing Values: " + str(df_bureau.isna().sum().sum()))
print("Data Frame Shape: " + str(df_bureau.shape))
print("Number of Missing Values by Feature: " + str(df_bureau.isna().sum()))
print("Data Types:",df_bureau.dtypes)
print("Data Frame: Data Types", df_bureau.dtypes.value_counts())
print("Summary", df_bureau.describe())
print("Correlation Statistic", df_bureau.corr())
print("Summary Information:", df_bureau.info())

```

```

Number of Rows: 1716428
Number of Columns: 17
Number of Total Missing Values: 3939947
Data Frame Shape: (1716428, 17)
Number of Missing Values by Feature: SK_ID_CURR      0
SK_ID_BUREAU      0
CREDIT_ACTIVE      0
CREDIT_CURRENCY      0
DAYS_CREDIT      0
CREDIT_DAY_OVERDUE      0
DAYS_CREDIT_ENDDATE      105553
DAYS_ENDDATE_FACT      633653
AMT_CREDIT_MAX_OVERDUE      1124488
CNT_CREDIT_PROLONG      0
AMT_CREDIT_SUM      13
AMT_CREDIT_SUM_DEBT      257669
AMT_CREDIT_SUM_LIMIT      591780
AMT_CREDIT_SUM_OVERDUE      0
CREDIT_TYPE      0
DAYS_CREDIT_UPDATE      0
AMT_ANNUITY      1226791
dtype: int64
Data Types: SK_ID_CURR      int64
SK_ID_BUREAU      int64
CREDIT_ACTIVE      object
CREDIT_CURRENCY      object
DAYS_CREDIT      int64
CREDIT_DAY_OVERDUE      int64
DAYS_CREDIT_ENDDATE      float64
DAYS_ENDDATE_FACT      float64
AMT_CREDIT_MAX_OVERDUE      float64
CNT_CREDIT_PROLONG      int64
AMT_CREDIT_SUM      float64
AMT_CREDIT_SUM_DEBT      float64
AMT_CREDIT_SUM_LIMIT      float64
AMT_CREDIT_SUM_OVERDUE      float64
CREDIT_TYPE      object
DAYS_CREDIT_UPDATE      int64
AMT_ANNUITY      float64
dtype: object
Data Frame: Data Types float64      8
int64      6
object      3
dtype: int64
Summary
count  1.716428e+06  1.716428e+06  1.716428e+06  1.716428e+06  \
mean    2.782149e+05  5.924434e+06 -1.142108e+03    8.181666e-01
std     1.029386e+05  5.322657e+05  7.951649e+02    3.654443e+01
min     1.000010e+05  5.000000e+06 -2.922000e+03    0.000000e+00
25%     1.888668e+05  5.463954e+06 -1.666000e+03    0.000000e+00
50%     2.780550e+05  5.926304e+06 -9.870000e+02    0.000000e+00
75%     3.674260e+05  6.385681e+06 -4.740000e+02    0.000000e+00
max     4.562550e+05  6.843457e+06  0.000000e+00    2.792000e+03

DAYS_CREDIT_ENDDATE  DAYS_ENDDATE_FACT  AMT_CREDIT_MAX_OVERDUE  \
count    1.610875e+06    1.082775e+06    5.919400e+05
mean     5.105174e+02   -1.017437e+03    3.825418e+03
std      4.994220e+03    7.140106e+02    2.060316e+05

```

✓ **File:** bureau_balance.csv


```

print("Number of Rows: " + str(df_bureau_balance.shape[0]))
print("Number of Columns: " + str(df_bureau_balance.shape[1]))
print("Number of Total Missing Values: " + str(df_bureau_balance.isna().sum().sum()))
print("Data Frame Shape: " + str(df_bureau_balance.shape))
print("Number of Missing Values by Feature: " + str(df_bureau_balance.isna().sum()))
print("Data Types:",df_bureau_balance.dtypes)
print("Data Frame: Data Types", df_bureau_balance.dtypes.value_counts())
print("Summary", df_bureau_balance.describe())
print("Correlation Statistic", df_bureau_balance.corr())
print("Summary Information:", df_bureau_balance.info())

```

Number of Rows: 27299925
 Number of Columns: 3
 Number of Total Missing Values: 0
 Data Frame Shape: (27299925, 3)
 Number of Missing Values by Feature: SK_ID_BUREAU 0
 MONTHS_BALANCE 0
 STATUS 0
 dtype: int64
 Data Types: SK_ID_BUREAU int64
 MONTHS_BALANCE int64
 STATUS object
 dtype: object
 Data Frame: Data Types int64 2
 object 1
 dtype: int64
 Summary

	SK_ID_BUREAU	MONTHS_BALANCE
count	2.729992e+07	2.729992e+07
mean	6.036297e+06	-3.074169e+01
std	4.923489e+05	2.386451e+01
min	5.001709e+06	-9.600000e+01
25%	5.730933e+06	-4.600000e+01
50%	6.070821e+06	-2.500000e+01
75%	6.431951e+06	-1.100000e+01
max	6.842888e+06	0.000000e+00

<ipython-input-31-1e0b9f61469f>:9: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a fu
 print("Correlation Statistic", df_bureau_balance.corr())
 Correlation Statistic

	SK_ID_BUREAU	MONTHS_BALANCE
SK_ID_BUREAU	1.000000	0.011873
MONTHS_BALANCE	0.011873	1.000000

<class 'pandas.core.frame.DataFrame'>
 RangeIndex: 27299925 entries, 0 to 27299924
 Data columns (total 3 columns):
 # Column Dtype

 0 SK_ID_BUREAU int64
 1 MONTHS_BALANCE int64
 2 STATUS object
 dtypes: int64(2), object(1)
 memory usage: 624.8+ MB
 Summary Information: None

✓ File: pos_cash_balance.csv

```

print("Number of Rows: " + str(df_pos_cash_balance.shape[0]))
print("Number of Columns: " + str(df_pos_cash_balance.shape[1]))
print("Number of Total Missing Values: " + str(df_pos_cash_balance.isna().sum().sum()))
print("Data Frame Shape: " + str(df_pos_cash_balance.shape))
print("Number of Missing Values by Feature: " + str(df_pos_cash_balance.isna().sum()))
print("Data Types:",df_pos_cash_balance.dtypes)
print("Data Frame: Data Types", df_pos_cash_balance.dtypes.value_counts())
print("Summary", df_pos_cash_balance.describe())
print("Correlation Statistic", df_pos_cash_balance.corr())
print("Summary Information:", df_pos_cash_balance.info())

```

```

CNT_INSTALMENT_FUTURE    SK_DPD    SK_DPD_DEF
count      9.975271e+06  1.000136e+07  1.000136e+07
mean       1.048384e+01  1.160693e+01  6.544684e-01
std        1.110906e+01  1.327140e+02  3.276249e+01
min        0.000000e+00  0.000000e+00  0.000000e+00
25%        3.000000e+00  0.000000e+00  0.000000e+00
50%        7.000000e+00  0.000000e+00  0.000000e+00
75%        1.400000e+01  0.000000e+00  0.000000e+00
max         8.500000e+01  4.231000e+03  3.595000e+03

```

```

<ipython-input-32-cb2b07444f43>:9: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future
print("Correlation Statistic", df_pos_cash_balance.corr())

```

```

Correlation Statistic      SK_ID_PREV  SK_ID_CURR  MONTHS_BALANCE  CNT_INSTALMENT  \
SK_ID_PREV      1.000000    -0.000336      0.001835      0.003820
SK_ID_CURR      -0.000336      1.000000      0.000404      0.000144
MONTHS_BALANCE      0.001835      0.000404      1.000000      0.336163
CNT_INSTALMENT      0.003820      0.000144      0.336163      1.000000
CNT_INSTALMENT_FUTURE  0.003679    -0.000559      0.271595      0.871276
SK_DPD           -0.000487      0.003118     -0.018939     -0.060803
SK_DPD_DEF       0.004848      0.001948     -0.000381     -0.014154

```

```

CNT_INSTALMENT_FUTURE  SK_DPD  SK_DPD_DEF
SK_ID_PREV      0.003679    -0.000487      0.004848
SK_ID_CURR     -0.000559      0.003118      0.001948
MONTHS_BALANCE      0.271595    -0.018939     -0.000381
CNT_INSTALMENT      0.871276    -0.060803     -0.014154
CNT_INSTALMENT_FUTURE  1.000000    -0.082004     -0.017436
SK_DPD           -0.082004      1.000000      0.245782
SK_DPD_DEF       -0.017436      0.245782      1.000000

```

```

<class 'pandas.core.frame.DataFrame'>

```

```

RangeIndex: 10001358 entries, 0 to 10001357

```

```

Data columns (total 8 columns):

```

```

#   Column      Dtype
---  -----  ---

```

```

0   SK_ID_PREV      int64
1   SK_ID_CURR      int64
2   MONTHS_BALANCE  int64
3   CNT_INSTALMENT  float64
4   CNT_INSTALMENT_FUTURE  float64
5   NAME_CONTRACT_STATUS  object
6   SK_DPD          int64
7   SK_DPD_DEF      int64

```

```

dtypes: float64(2), int64(5), object(1)

```

```

memory usage: 610.4+ MB

```

```

Summary Information: None

```

File: credit_card_balance.csv

```

print("Number of Rows: " + str(df_credit_card_balance.shape[0]))
print("Number of Columns: " + str(df_credit_card_balance.shape[1]))
print("Number of Total Missing Values: " + str(df_credit_card_balance.isna().sum().sum()))
print("Data Frame Shape: " + str(df_credit_card_balance.shape))
print("Number of Missing Values by Feature: " + str(df_credit_card_balance.isna().sum()))
print("Data Types:", df_credit_card_balance.dtypes)
print("Data Frame: Data Types", df_credit_card_balance.dtypes.value_counts())
print("Summary", df_credit_card_balance.describe())
print("Correlation Statistic", df_credit_card_balance.corr())
print("Summary Information:", df_credit_card_balance.info())

```

SK_DP
SK_DPD_DEF

0.002156 0.218950 1.000000

```
[22 rows x 22 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3840312 entries, 0 to 3840311
Data columns (total 23 columns):
#   Column                                Dtype
---  -
0   SK_ID_PREV                           int64
1   SK_ID_CURR                           int64
2   MONTHS_BALANCE                       int64
3   AMT_BALANCE                          float64
4   AMT_CREDIT_LIMIT_ACTUAL              int64
5   AMT_DRAWINGS_ATM_CURRENT             float64
6   AMT_DRAWINGS_CURRENT                 float64
7   AMT_DRAWINGS_OTHER_CURRENT           float64
8   AMT_DRAWINGS_POS_CURRENT             float64
9   AMT_INST_MIN_REGULARITY              float64
10  AMT_PAYMENT_CURRENT                  float64
11  AMT_PAYMENT_TOTAL_CURRENT             float64
12  AMT_RECEIVABLE_PRINCIPAL              float64
13  AMT_RECIVABLE                         float64
14  AMT_TOTAL_RECEIVABLE                  float64
15  CNT_DRAWINGS_ATM_CURRENT              float64
16  CNT_DRAWINGS_CURRENT                  int64
17  CNT_DRAWINGS_OTHER_CURRENT            float64
18  CNT_DRAWINGS_POS_CURRENT              float64
19  CNT_INSTALLMENT_MATURE_CUM            float64
20  NAME_CONTRACT_STATUS                  object
21  SK_DPD                               int64
22  SK_DPD_DEF                           int64
dtypes: float64(15), int64(7), object(1)
memory usage: 673.9+ MB
Summary Information: None
```

✓ File: previous_application.csv

```
print("Number of Rows: " + str(df_previous_application.shape[0]))
print("Number of Columns: " + str(df_previous_application.shape[1]))
print("Number of Total Missing Values: " + str(df_previous_application.isna().sum().sum()))
print("Data Frame Shape: " + str(df_previous_application.shape))
print("Number of Missing Values by Feature: " + str(df_previous_application.isna().sum()))
print("Data Types:", df_previous_application.dtypes)
print("Data Frame: Data Types", df_previous_application.dtypes.value_counts())
print("Summary", df_previous_application.describe())
print("Correlation Statistic", df_previous_application.corr())
print("Summary Information:", df_previous_application.info())
```

```
Number of Rows: 1670214
Number of Columns: 37
Number of Total Missing Values: 11109336
Data Frame Shape: (1670214, 37)
Number of Missing Values by Feature: SK_ID_PREV 0
SK_ID_CURR 0
NAME_CONTRACT_TYPE 0
AMT_ANNUITY 372235
AMT_APPLICATION 0
AMT_CREDIT 1
AMT_DOWN_PAYMENT 895844
AMT_GOODS_PRICE 385515
WEEKDAY_APPR_PROCESS_START 0
HOUR_APPR_PROCESS_START 0
FLAG_LAST_APPL_PER_CONTRACT 0
NFLAG_LAST_APPL_IN_DAY 0
RATE_DOWN_PAYMENT 895844
RATE_INTEREST_PRIMARY 1664263
RATE_INTEREST_PRIVILEGED 1664263
NAME_CASH_LOAN_PURPOSE 0
NAME_CONTRACT_STATUS 0
DAYS_DECISION 0
NAME_PAYMENT_TYPE 0
CODE_REJECT_REASON 0
NAME_TYPE_SUITE 820405
NAME_CLIENT_TYPE 0
NAME_GOODS_CATEGORY 0
NAME_PORTFOLIO 0
NAME_PRODUCT_TYPE 0
CHANNEL_TYPE 0
SELLERPLACE_AREA 0
NAME_SELLER_INDUSTRY 0
```

```

CNT_PAYMENT          372230
NAME_YIELD_GROUP      0
PRODUCT_COMBINATION   346
DAYS_FIRST_DRAWING    673065
DAYS_FIRST_DUE        673065
DAYS_LAST_DUE_1ST_VERSION 673065
DAYS_LAST_DUE         673065
DAYS_TERMINATION      673065
NFLAG_INSURED_ON_APPROVAL 673065
dtype: int64
Data Types: SK_ID_PREV          int64
SK_ID_CURR          int64
NAME_CONTRACT_TYPE   object
AMT_ANNUITY          float64
AMT_APPLICATION      float64
AMT_CREDIT           float64
AMT_DOWN_PAYMENT     float64
AMT_GOODS_PRICE      float64
WEEKDAY_APPR_PROCESS_START object
HOUR_APPR_PROCESS_START int64
FLAG_LAST_APPL_PER_CONTRACT object
NFLAG_LAST_APPL_IN_DAY int64
RATE_DOWN_PAYMENT    float64
RATE_INTEREST_PRIMARY float64
RATE_INTEREST_PRIVILEGED float64
NAME_CASH_LOAN_PURPOSE object

```

✓ File: installments_payments.csv

```

print("Number of Rows: " + str(df_installments_payments.shape[0]))
print("Number of Columns: " + str(df_installments_payments.shape[1]))
print("Number of Total Missing Values: " + str(df_installments_payments.isna().sum().sum()))
print("Data Frame Shape: " + str(df_installments_payments.shape))
print("Number of Missing Values by Feature: " + str(df_installments_payments.isna().sum()))
print("Data Types:", df_installments_payments.dtypes)
print("Data Frame: Data Types", df_installments_payments.dtypes.value_counts())
print("Summary", df_installments_payments.describe())
print("Correlation Statistic", df_installments_payments.corr())
print("Summary Information:", df_installments_payments.info())

```

```

      AMT_PAYMENT      0.120000
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13605401 entries, 0 to 13605400
Data columns (total 8 columns):
 #   Column              Dtype
---  -
 0   SK_ID_PREV          int64
 1   SK_ID_CURR          int64
 2   NUM_INSTALLMENT_VERSION float64
 3   NUM_INSTALLMENT_NUMBER int64
 4   DAYS_INSTALLMENT    float64
 5   DAYS_ENTRY_PAYMENT  float64
 6   AMT_INSTALLMENT     float64
 7   AMT_PAYMENT         float64
dtypes: float64(5), int64(3)
memory usage: 830.4 MB
Summary Information: None

```

✓ Visual Exploratory Data Analysis

✓ File: application_train.csv

```

import matplotlib.pyplot as plt
import seaborn as sns

```

✓ Feature: Target

Feature description: Target variable (1 - client with payment difficulties: he/she had late payment more than X days on at least one of the first Y installments of the loan in our sample, 0 - all other cases)

```

import matplotlib.patches as mpatches

fig, ax = plt.subplots(figsize=(8, 6))

barplot = df_application_train['TARGET'].value_counts().plot(kind='bar', color=['cyan', 'lightblue'])


barplot.set_title("Count of Target Feature")
barplot.set_ylabel("Count")
barplot.set_xlabel("Target Values")

legend_labels = ["People didn't face difficulties", "People faced difficulties"]
legend_patches = [mpatches.Patch(color=color, label=label) for color, label in zip(['cyan', 'lightblue'], legend_labels)]
plt.legend(handles=legend_patches, loc='upper right')

for i, count in enumerate(df_application_train['TARGET'].value_counts()):
    plt.text(i, count + 500, str(count), ha='center', va='bottom')

plt.show()

```




We can see from the above that:

- People who faced difficulties in repaying the loan sum up to **[Class 1]: 24825**
- People who didn't face any difficulties repaying the loan sum up to **[Class 0]: 282686**

```
df_calculated_age = df_application_train['DAYS_BIRTH']//365
fig, ax = plt.subplots(figsize=(8, 6))

ax.set_title('Age Distribution')
ax.set_ylabel('Count')
ax.set_xlabel('Age')
ax.hist(df_calculated_age, bins=25, color="green", edgecolor="black")

plt.show()
```




Analysis

```
fig, ax = plt.subplots(figsize=(8, 6))

ax.set_title('Gender Distribution')
ax.set_ylabel('Count')
ax.set_xlabel('Gender')
ax.hist(df_application_train['CODE_GENDER'], color="lightblue", edgecolor="black")

plt.show()
```

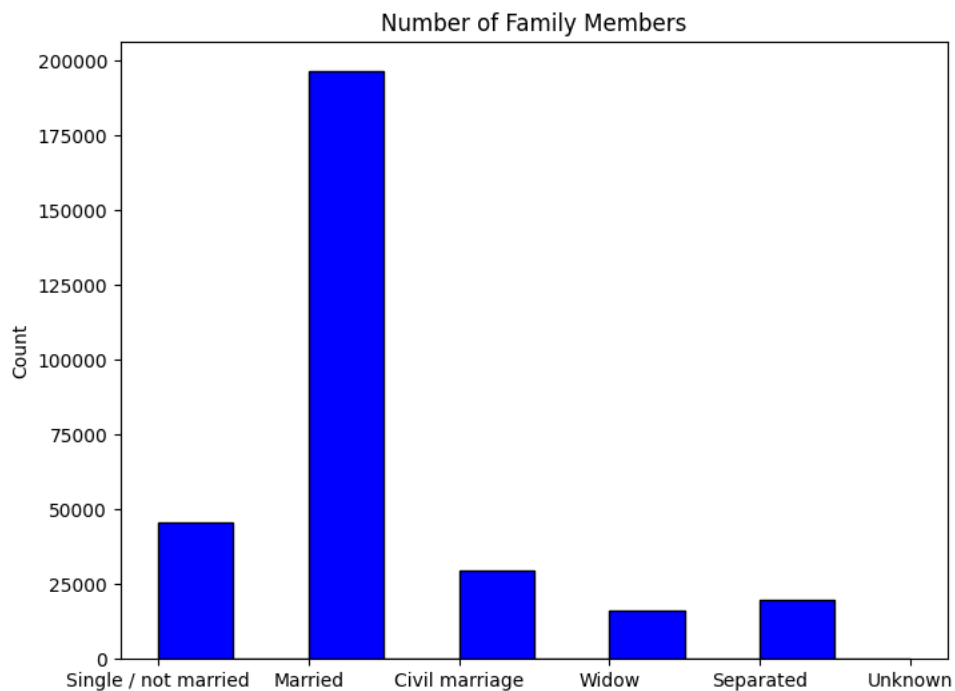


```
import seaborn as sns
import matplotlib.pyplot as plt

fig, ax = plt.subplots(figsize=(8, 6))

ax.set_title('Number of Family Members')
ax.set_ylabel('Count')
ax.set_xlabel('')
ax.hist(df_application_train['NAME_FAMILY_STATUS'], color="blue", edgecolor="black")

plt.show()
```




```
fig, ax = plt.subplots(figsize=(8, 6))
ax.set_title('Number of Family Members')
ax.set_ylabel('Count')
ax.set_xlabel('Number of Family Members')

sns.countplot(ax=ax, data=df_application_train, palette="pastel", x="CNT_FAM_MEMBERS")

ax.grid(axis='y', linestyle='--', alpha=0.7)

plt.show()
```




```
fig, ax = plt.subplots(figsize=(8, 6))
ax.set_title('Distribution for number of children')
ax.set_ylabel('Count')
ax.set_xlabel('Number of Children')

sns.countplot(ax=ax, data=df_application_train, palette="pink", x="CNT_CHILDREN")

ax.grid(axis='y', linestyle='--', alpha=0.7)

plt.show()
```



```
education_type_counts = df_application_train['NAME_EDUCATION_TYPE'].value_counts()

print(education_type_counts)
```


```
Secondary / secondary special    218391
Higher education                 74863
Incomplete higher                10277
Lower secondary                 3816
Academic degree                 164
Name: NAME_EDUCATION_TYPE, dtype: int64
```



```
fig, ax = plt.subplots(figsize=(14, 10))

ax.set_title('Income Distribution')
ax.set_ylabel('Count')
ax.set_xlabel('Income')
ax.hist(df_application_train['NAME_INCOME_TYPE'], bins=25, color="green", edgecolor="black")
ax.grid(axis='y', linestyle='--', alpha=0.7)

plt.show()
```



✓ Let's check correlation between features and target variable.

```
import seaborn as sns
import matplotlib.pyplot as plt


features = df_application_train[['NAME_FAMILY_STATUS', 'TARGET']]

# Compute the correlation matrix for the two features
corr_matrix = features.corr()

# Create a heatmap for the correlation matrix
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Heatmap for Correlation between NAME_FAMILY_STATUS and TARGET')
plt.show()
```

```
<ipython-input-45-858e97cceb8f>:8: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a fu
corr_matrix = features.corr()
```

Heatmap for Correlation between NAME_FAMILY_STATUS and TARGET



```
spearman_correlation = df_application_train['NAME_FAMILY_STATUS'].corr(df_application_train['TARGET'], method='spearman')
print(f"The Spearman correlation between NAME_FAMILY_STATUS and TARGET is: {spearman_correlation}")
```

The Spearman correlation between NAME_FAMILY_STATUS and TARGET is: -0.001815167766133806

```
df_application_train.head()
```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CRE
0	100002	1	Cash loans	M	N	Y	0	202500.0	4065
1	100003	0	Cash loans	F	N	N	0	270000.0	12935
2	100004	0	Revolving loans	M	Y	Y	0	67500.0	1350
3	100006	0	Cash loans	F	N	Y	0	135000.0	3126
4	100007	0	Cash loans	M	N	Y	0	121500.0	5130

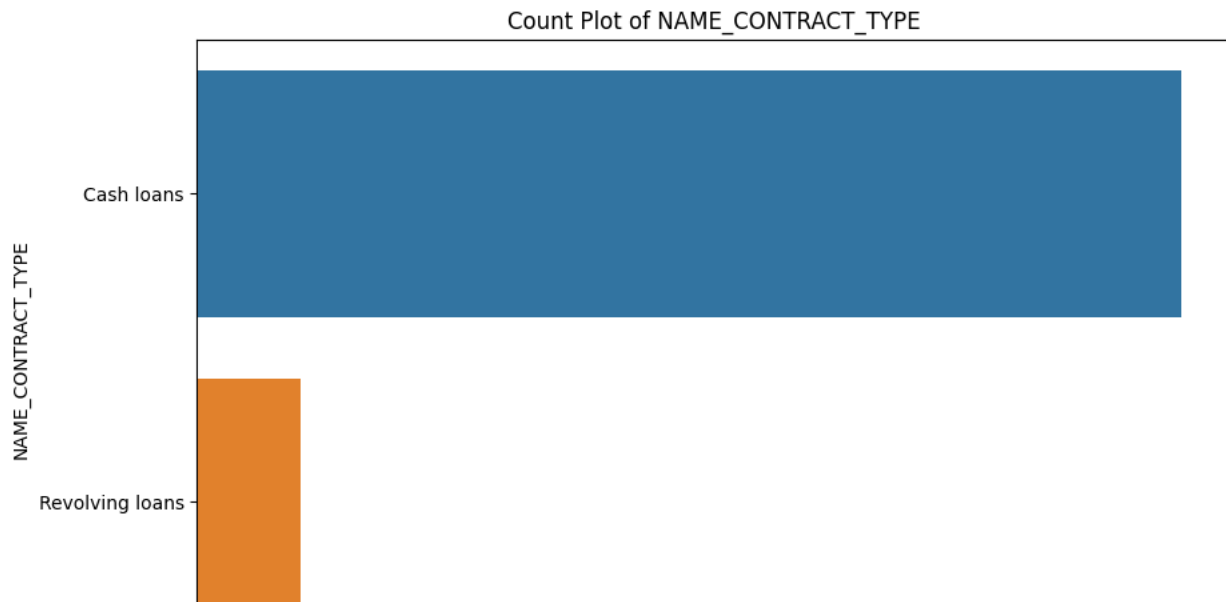
5 rows x 122 columns

```
column_names=df_application_train.columns.tolist()
print(column_names)
```

```
['SK_ID_CURR', 'TARGET', 'NAME_CONTRACT_TYPE', 'CODE_GENDER', 'FLAG_OWN_CAR', 'FLAG_OWN_REALTY', 'CNT_CHILDREN', 'AMT_INCOME'
```


```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Count Plot for a categorical variable
plt.figure(figsize=(10, 6))
sns.countplot(data=df_application_train, y='NAME_CONTRACT_TYPE')
plt.title('Count Plot of NAME_CONTRACT_TYPE')
plt.xlabel('Count')
plt.ylabel('NAME_CONTRACT_TYPE')
plt.show()
```



```
import seaborn as sns
import matplotlib.pyplot as plt
```

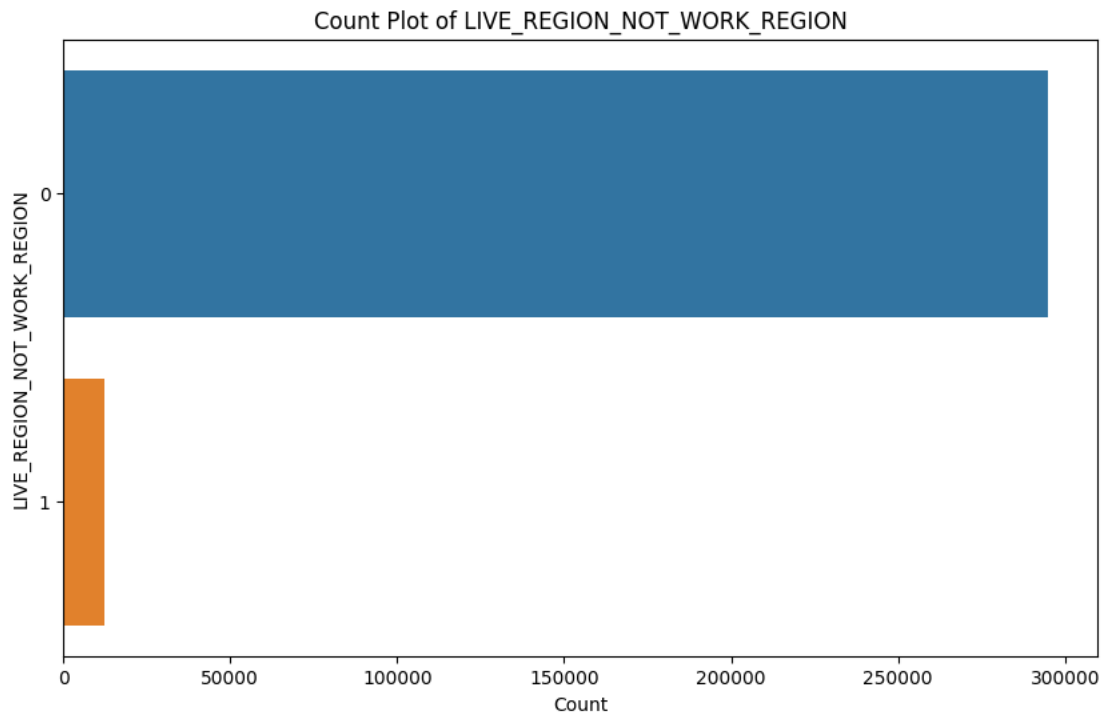
```
# Count Plot for a categorical variable
plt.figure(figsize=(10, 6))
sns.countplot(data=df_application_train, y='OCCUPATION_TYPE')
plt.title('Count Plot of OCCUPATION_TYPE')
plt.xlabel('Count')
plt.ylabel('OCCUPATION_TYPE')
plt.show()
```



- ✓ This categorical feature tend to display which occupations are most likely to do loan application. Previously.


```
import seaborn as sns
import matplotlib.pyplot as plt

# Count Plot for a categorical variable
plt.figure(figsize=(10, 6))
sns.countplot(data=df_application_train, y='LIVE_REGION_NOT_WORK_REGION')
plt.title('Count Plot of LIVE_REGION_NOT_WORK_REGION')
plt.xlabel('Count')
plt.ylabel('LIVE_REGION_NOT_WORK_REGION')
plt.show()
```



```
import matplotlib.pyplot as plt
import seaborn as sns

# Histogram of DAYS_FIRST_DUE
plt.figure(figsize=(10, 6))
sns.histplot(df_application_train['TOTALAREA_MODE'], bins=30, kde=True)
plt.title('Distribution of TOTALAREA_MODE')
plt.xlabel('TOTALAREA_MODE')
plt.ylabel('Frequency')
plt.show()
```



Left skewed graph, shows us that total area owned by the applicants which can be used as credibility to check whether to grant loan or not

Handling missing values in application_train

```
missing_values = df_application_train.isnull()
print('Missing Values in entire dataframe',missing_values)
```

	Missing Values in entire dataframe	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
...
307506	False	False	False	False	False	False
307507	False	False	False	False	False	False
307508	False	False	False	False	False	False
307509	False	False	False	False	False	False
307510	False	False	False	False	False	False

	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
307506	False	False	False	False
307507	False	False	False	False
307508	False	False	False	False
307509	False	False	False	False
307510	False	False	False	False

	AMT_ANNUITY	FLAG_DOCUMENT_18	FLAG_DOCUMENT_19
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...
307506	False	False	False
307507	False	False	False
307508	False	False	False
307509	False	False	False
307510	False	False	False

	FLAG_DOCUMENT_20	FLAG_DOCUMENT_21	AMT_REQ_CREDIT_BUREAU_HOUR
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	True
4	False	False	False
...
307506	False	False	True
307507	False	False	True
307508	False	False	False
307509	False	False	False
307510	False	False	False

	AMT_REQ_CREDIT_BUREAU_DAY	AMT_REQ_CREDIT_BUREAU_WEEK
0	False	False
1	False	False
2	False	False
3	True	True
4	False	False

```
missing_values_count = df_application_train.isnull().sum()
print(missing_values_count)
```

```
SK_ID_CURR      0
TARGET          0
NAME_CONTRACT_TYPE 0
CODE_GENDER     0
FLAG_OWN_CAR    0
```

```
...
AMT_REQ_CREDIT_BUREAU_DAY 41519
AMT_REQ_CREDIT_BUREAU_WEEK 41519
AMT_REQ_CREDIT_BUREAU_MON 41519
AMT_REQ_CREDIT_BUREAU_QRT 41519
AMT_REQ_CREDIT_BUREAU_YEAR 41519
Length: 122, dtype: int64
```

```
missing_values_percentage = (df_application_train.isnull().sum() / len(df_application_train)) * 100
###missing values percent per column
print(missing_values_percentage)
```

```
SK_ID_CURR      0.000000
TARGET          0.000000
NAME_CONTRACT_TYPE 0.000000
CODE_GENDER     0.000000
FLAG_OWN_CAR    0.000000
...
AMT_REQ_CREDIT_BUREAU_DAY 13.501631
AMT_REQ_CREDIT_BUREAU_WEEK 13.501631
AMT_REQ_CREDIT_BUREAU_MON 13.501631
AMT_REQ_CREDIT_BUREAU_QRT 13.501631
AMT_REQ_CREDIT_BUREAU_YEAR 13.501631
Length: 122, dtype: float64
```

```
missing_values_rows = df_application_train[df_application_train.isnull().any(axis=1)]
###filtering dataframe to show only missing values
print(missing_values_rows)
```

```
SK_ID_CURR  TARGET  NAME_CONTRACT_TYPE  CODE_GENDER  FLAG_OWN_CAR  \
0          100002      1      Cash loans           M           N
1          100003      0      Cash loans           F           N
2          100004      0  Revolving loans           M           Y
3          100006      0      Cash loans           F           N
4          100007      0      Cash loans           M           N
...          ...      ...      ...      ...      ...
307506      456251      0      Cash loans           M           N
307507      456252      0      Cash loans           F           N
307508      456253      0      Cash loans           F           N
307509      456254      1      Cash loans           F           N
307510      456255      0      Cash loans           F           N
```

```
FLAG_OWN_REALTY  CNT_CHILDREN  AMT_INCOME_TOTAL  AMT_CREDIT  \
0                Y            0      202500.0      406597.5
1                N            0      270000.0     1293502.5
2                Y            0      67500.0     135000.0
3                Y            0     135000.0     312682.5
4                Y            0     121500.0     513000.0
...          ...      ...      ...      ...
307506          N            0     157500.0     254700.0
307507          Y            0      72000.0     269550.0
307508          Y            0     153000.0     677664.0
307509          Y            0     171000.0     370107.0
307510          N            0     157500.0     675000.0
```

```
AMT_ANNUITY  ...  FLAG_DOCUMENT_18  FLAG_DOCUMENT_19  FLAG_DOCUMENT_20  \
0      24700.5  ...            0            0            0
1      35698.5  ...            0            0            0
2       6750.0  ...            0            0            0
3      29686.5  ...            0            0            0
4      21865.5  ...            0            0            0
...          ...      ...      ...      ...
307506      27558.0  ...            0            0            0
307507      12001.5  ...            0            0            0
307508      29979.0  ...            0            0            0
307509      20205.0  ...            0            0            0
307510      49117.5  ...            0            0            0
```

```
FLAG_DOCUMENT_21  AMT_REQ_CREDIT_BUREAU_HOUR  AMT_REQ_CREDIT_BUREAU_DAY  \
0                0                0.0                0.0
1                0                0.0                0.0
2                0                0.0                0.0
3                0                NaN                NaN
4                0                0.0                0.0
...          ...      ...      ...
307506          0                NaN                NaN
307507          0                NaN                NaN
```

```

307508      0      1.0      0.0
307509      0      0.0      0.0
307510      0      0.0      0.0

```

```

      AMT_REQ_CREDIT_BUREAU_WEEK  AMT_REQ_CREDIT_BUREAU_MON  \
0      0.0      0.0
1      0.0      0.0
2      0.0      0.0
3      NaN      NaN
4      0.0      0.0

```

```


import seaborn as sns
import matplotlib.pyplot as plt

```

```

plt.figure(figsize=(10, 6))
sns.heatmap(df_application_train.isnull(), cbar=False, cmap='viridis')
plt.show()

```



✓ Heatmap of missing values in application_train

```

missing_values = df_application_train['DAYS_BIRTH'].isnull()
print(missing_values)

```

```

df_application_train.dropna(inplace=True)

```

```
correlation_data = df_application_train[["TARGET", "CNT_CHILDREN", "CNT_FAM_MEMBERS", "DAYS_BIRTH"]]
correlation_data["DAYS_BIRTH"] = abs(correlation_data["DAYS_BIRTH"])
correlation_data = correlation_data.corr()
```

```
fig, ax = plt.subplots(figsize=(10, 8))
cax = ax.matshow(correlation_data, cmap="magma", vmin=-1.0, vmax=1.0)
fig.colorbar(cax)
```

```
ticks = list(range(len(correlation_data.columns)))
```

```
ax.set_xticks(ticks)
```

```
ax.set_yticks(ticks)
```

```
ax.set_xticklabels(correlation_data.columns, rotation=45, ha="left")
```

```
ax.set_yticklabels(correlation_data.columns)
```

```
for i in range(len(correlation_data.columns)):
```

```
    for j in range(len(correlation_data.columns)):
```

```
        text = ax.text(j, i, f"{correlation_data.iloc[i, j]:.2f}", ha="center", va="center", color="black")
```


```
plt.title("Correlation Heatmap: application_train")
```

```
plt.show()
```

<ipython-input-60-1d187026b71b>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view
correlation_data["DAYS_BIRTH"] = abs(correlation_data["DAYS_BIRTH"])

Correlation Heatmap: application_train



```
datacorr = df_application_train.corrwith(df_application_train['TARGET'])
datacorr
```

<ipython-input-61-25930280b08d>:1: FutureWarning: The default value of numeric_only in DataFrame.corrwith is deprecated. In
datacorr = df_application_train.corrwith(df_application_train['TARGET'])
SK_ID_CURR 0.015474
TARGET 1.000000


```

CNT_CHILDREN      -0.019993
AMT_INCOME_TOTAL  -0.039762
AMT_CREDIT        -0.014634
...
AMT_REQ_CREDIT_BUREAU_DAY    0.014616
AMT_REQ_CREDIT_BUREAU_WEEK  0.015000
AMT_REQ_CREDIT_BUREAU_MON   -0.004202
AMT_REQ_CREDIT_BUREAU_QRT   0.016465
AMT_REQ_CREDIT_BUREAU_YEAR  0.033832
Length: 106, dtype: float64

```

```

import seaborn as sns
import matplotlib.pyplot as plt


app_sorted = datacorr.sort_values()
plt.figure(figsize=(70, 6))
sns.barplot(x=app_sorted.index, y=app_sorted.values, palette='viridis')

plt.title('Correlation with "target_column" (Sorted)')
plt.xlabel('Columns')
plt.ylabel('Correlation')

plt.xticks(rotation=45)
plt.tight_layout()

plt.show()

```



```

correlation_data = df_application_train[['TARGET', 'EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'REGION_RATING_CLIENT_W_CITY', 'R
correlation_data["DAYS_BIRTH"] = abs(correlation_data["DAYS_BIRTH"])
correlation_data = correlation_data.corr()

fig, ax = plt.subplots(figsize=(10, 8))
cax = ax.matshow(correlation_data, cmap="magma", vmin=-1.0, vmax=1.0)
fig.colorbar(cax)

ticks = list(range(len(correlation_data.columns)))
ax.set_xticks(ticks)
ax.set_yticks(ticks)
ax.set_xticklabels(correlation_data.columns, rotation=45, ha="left")
ax.set_yticklabels(correlation_data.columns)

for i in range(len(correlation_data.columns)):
    for j in range(len(correlation_data.columns)):
        text = ax.text(j, i, f"{correlation_data.iloc[i, j]:.2f}", ha="center", va="center", color="black")

plt.title("Correlation Heatmap: \napplcation_train - bEST FEATURES")
plt.show()

```


12/5/23, 10:11 PM

FP_GroupN_HCDR_16_Phase 4.ipynb - Colaboratory


<ipython-input-63-4111a7d6d40f>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view
correlation_data["DAYS_BIRTH"] = abs(correlation_data["DAYS_BIRTH"])

Correlation Heatmap:
application_train - BEST FEATURES



Implementation of Neural Networks



```

import torch
import torch.utils.data
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, roc_auc_score, roc_curve, auc, f1_score
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
import matplotlib.pyplot as plt
import datetime
import random
import string

from torch.utils.data import Dataset, TensorDataset, DataLoader
from sklearn.feature_selection import VarianceThreshold

torch.manual_seed(0)
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

# load data
hcdr_application = pd.read_csv("clean_final.csv")

X = hcdr_application.drop('TARGET', axis = 1)
y = hcdr_application.TARGET
print("Shapes:", X.shape, y.shape)

# train test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, shuffle = True)
X_train, X_validation, y_train, y_validation = train_test_split(X_train, y_train, test_size=0.15, random_state=42, shuffle=True)

y_train = y_train.to_numpy()
y_validation = y_validation.to_numpy()
y_test = y_test.to_numpy()

# convert numpy arrays to tensors
X_train_tensor = torch.from_numpy(np.array(X_train))
X_valid_tensor = torch.from_numpy(np.array(X_validation))
X_test_tensor = torch.from_numpy(np.array(X_test))
y_train_tensor = torch.from_numpy(y_train.astype('int'))
y_valid_tensor = torch.from_numpy(y_validation.astype('int'))
y_test_tensor = torch.from_numpy(y_test.astype('int'))

# create TensorDataset in PyTorch
hcdr_train = torch.utils.data.TensorDataset(X_train_tensor, y_train_tensor)
hcdr_valid = torch.utils.data.TensorDataset(X_valid_tensor, y_valid_tensor)
hcdr_test = torch.utils.data.TensorDataset(X_test_tensor, y_test_tensor)

train_batch_size = 32
valid_test_batch_size = 16
trainloader_hcdr = torch.utils.data.DataLoader(hcdr_train, batch_size=train_batch_size, shuffle=True, num_workers=2)
validloader_hcdr = torch.utils.data.DataLoader(hcdr_valid, batch_size=valid_test_batch_size, shuffle=True, num_workers=2)
testloader_hcdr = torch.utils.data.DataLoader(hcdr_test, batch_size=valid_test_batch_size, shuffle=True, num_workers=2)

def run_hcdr_model(
    hidden_layer_neurons=[32, 16, 8],
    opt=optim.SGD,
    epochs=5,
    learning_rate=1e-3
):
    D_in = X_test.shape[1] # Input layer neurons depend on the input dataset shape
    D_out = 2 # Output layer neurons here, 2 classes: 0 and 1

```



```

str_neurons = [str(h) for h in hidden_layer_neurons]
arch_string = f"{D_in}-{ '-' .join(str_neurons) }-{D_out}"

layers = [
    torch.nn.Linear(D_in, hidden_layer_neurons[0]), # X.matmul(W1)
    nn.ReLU(), # ReLU( X.matmul(W1))
]

# Add hidden layers
for i in range(1, len(hidden_layer_neurons)):
    prev, curr = hidden_layer_neurons[i - 1], hidden_layer_neurons[i]
    layers.append(torch.nn.Linear(prev, curr))
    layers.append(nn.ReLU())
# layers.append(nn.Sigmoid())

# Add final layer

layers.append(nn.Linear(hidden_layer_neurons[-1], D_out)) # Relu( X.matmul(W1)).matmul(W2))

model = torch.nn.Sequential(*layers)

model.to(device)

loss_fn = nn.CrossEntropyLoss() #for classification
optimizer = opt(model.parameters(), lr=learning_rate)

#summary(model, (4, 20))
print('-'*50)
print('Model:')
print(model)
print('-'*50)

loss_history = []
acc_history = []
def train_epoch(epoch, model, loss_fn, opt, train_loader):
    running_loss = 0.0
    count = 0
    y_pred = []
    epoch_target = []

    for batch_id, data in enumerate(train_loader):
        inputs, target = data[0].to(device), data[1].to(device)
        # 1:zero the grad, 2:forward pass, 3:calculate loss, and 4:backprop!
        opt.zero_grad()
        preds = model(inputs.float())

        # compute loss and gradients
        loss = loss_fn(preds, target)

        loss.backward() #calculate nabla_w
        loss_history.append(loss.item())
        opt.step() #update W
        y_pred.extend(torch.argmax(preds, dim=1).tolist())
        epoch_target.extend(target.tolist())

        running_loss += loss.item()
        count += 1

    loss = np.round(running_loss/count, 3)

    #accuracy
    correct = (np.array(y_pred) == np.array(epoch_target))
    accuracy = correct.sum() / correct.size
    accuracy = np.round(accuracy, 3)
    return loss, accuracy

def evaluate_model(epoch, model, loss_fn, opt, data_loader, tag = "Test"):
    overall_loss = 0.0
    count = 0

```

```

y_pred = []
epoch_target = []
for i,data in enumerate(data_loader):
    inputs, target = data[0].to(device), data[1].to(device)
    preds = model(inputs.float())

    loss = loss_fn(preds, target)          # compute loss value

    overall_loss += (loss.item()) # compute total loss to save to logs
    y_pred.extend(torch.argmax(preds, dim=1).tolist())
    epoch_target.extend(target.tolist())
    count += 1

# compute mean loss
loss = np.round(overall_loss/count, 3)
#accuracy
correct = (np.array(y_pred) == np.array(epoch_target))
accuracy = correct.sum() / correct.size
accuracy = np.round(accuracy, 3)
return loss, accuracy

for epoch in range(epochs):
    # print(f"Epoch {epoch+1}")
    train_loss, train_accuracy = train_epoch(epoch, model, loss_fn, optimizer, trainloader_hcdr)
    valid_loss, valid_accuracy = evaluate_model(epoch, model, loss_fn, optimizer, validloader_hcdr, tag = "Validation")
    print(f"Epoch {epoch+1}: Train Accuracy: {train_accuracy}\t Validation Accuracy: {valid_accuracy}")
print("-"*50)
test_loss, test_accuracy = evaluate_model(epoch, model, loss_fn, opt, testloader_hcdr, tag="Test")

return arch_string, train_accuracy, valid_accuracy, test_accuracy, model

Shapes: (307511, 246) (307511,)

```

```

import pandas as pd
torch.manual_seed(0)

# hidden_layer_neurons = [32,16,8]
hidden_layer_neurons = [300,200,64,8]
opt = optim.Adam # optim.SGD, Optim.Adam, etc.
epochs = 3
learning_rate = 2e-3

arch_string, train_accuracy, valid_accuracy, test_accuracy,model = run_hcdr_model(
    hidden_layer_neurons,
    opt,
    epochs,
    learning_rate
)

try: Log
except : Log = pd.DataFrame(
    columns=[
        "Architecture string",
        "Optimizer",
        "Epochs",
        "Train accuracy",
        "Valid accuracy",
        "Test accuracy",
    ]
)

Log.loc[len(Log)] = [
    arch_string,
    f"{opt}",
    f"{epochs}",
    f"{train_accuracy * 100}%",
    f"{valid_accuracy * 100}%",
    f"{test_accuracy * 100}%",
]

Log

```

```

-----
Model:
Sequential(
  (0): Linear(in_features=246, out_features=300, bias=True)
  (1): ReLU()
  (2): Linear(in_features=300, out_features=200, bias=True)
  (3): ReLU()
  (4): Linear(in_features=200, out_features=64, bias=True)
  (5): ReLU()
  (6): Linear(in_features=64, out_features=8, bias=True)
  (7): ReLU()
  (8): Linear(in_features=8, out_features=2, bias=True)
)

```

```

-----
Epoch 1: Train Accuracy: 0.911   Validation Accuracy: 0.919
Epoch 2: Train Accuracy: 0.919   Validation Accuracy: 0.919
Epoch 3: Train Accuracy: 0.919   Validation Accuracy: 0.919
-----

```


	Architecture string	Optimizer	Epochs	Train accuracy	Valid accuracy	Test accuracy
0	246-32-16-8-2	<class 'torch.optim.adam.Adam'>	3	91.9%	91.9%	92.0%
1	246-300-200-64-8-2	<class 'torch.optim.adam.Adam'>	3	91.9%	91.9%	92.0%

```
with torch.no_grad():
    model.eval()
    outputs = model(X_test_tensor.float().cuda())

predicted_probs = outputs[:, 1].cpu().numpy()
predicted_probs[predicted_probs<0] = 0
predicted_probs[predicted_probs>0] = 1
true_labels = y_test_tensor.numpy()

# Calculate ROC curve and AUC
fpr, tpr, thresholds = roc_curve(true_labels, predicted_probs)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'AUC = {roc_auc:.2f}')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```




```

from sklearn.preprocessing import LabelEncoder
pd.set_option('display.float_format', lambda x: '%.5f' % x)
pd.set_option('mode.chained_assignment', None)

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
if torch.cuda.is_available():
    print("GPU - ", torch.cuda.get_device_name(0))

# SET HYPERPARAMETERS
test_size = 0.2
epochs = 12
batch_size = 320
learningrate= 0.000008
hp_emb_drop = 0.04
nnlayers = [800, 350]
hp_ps = [0.001, 0.01]

# LOAD DATA
application_train_df = pd.read_csv('dataset/application_train.csv').sample(frac = 1)
application_test_df = pd.read_csv('dataset/application_test.csv')
previous_application_df = pd.read_csv('dataset/previous_application.csv')

application_train_df['CSV_SOURCE'] = 'application_train.csv'
application_test_df['CSV_SOURCE'] = 'application_test.csv'
df = pd.concat([application_train_df, application_test_df])

# PREPARING previous_applications.csv
temp_previous_df = previous_application_df.groupby('SK_ID_CURR', as_index=False).agg({'NAME_CONTRACT_STATUS': lambda x: ','.join(x)})
temp_previous_df['has_only_approved'] = np.where(temp_previous_df['NAME_CONTRACT_STATUS'] == 'Approved', '1', '0')
temp_previous_df['has_been_rejected'] = np.where(temp_previous_df['NAME_CONTRACT_STATUS'].str.contains('Refused'), '1', '0')

# JOIN DATA
df = pd.merge(df, temp_previous_df, on='SK_ID_CURR', how='left')

# Feature engineering
# total_amt_req_credit_bureau
df['total_amt_req_credit_bureau'] = (
    df['AMT_REQ_CREDIT_BUREAU_YEAR'] * 1 +
    df['AMT_REQ_CREDIT_BUREAU_QRT'] * 2 +
    df['AMT_REQ_CREDIT_BUREAU_MON'] * 8 +
    df['AMT_REQ_CREDIT_BUREAU_WEEK'] * 16 +
    df['AMT_REQ_CREDIT_BUREAU_DAY'] * 32 +
    df['AMT_REQ_CREDIT_BUREAU_HOUR'] * 64)
df['total_amt_req_credit_bureau_isnull'] = np.where(df['total_amt_req_credit_bureau'].isnull(), '1', '0')
df['total_amt_req_credit_bureau'].fillna(0, inplace=True)

#has_job
df['has_job'] = np.where(df['NAME_INCOME_TYPE'].isin(['Pensioner', 'Student', 'Unemployed']), '1', '0')

# has_children
df['has_children'] = np.where(df['CNT_CHILDREN'] > 0, '1', '0')

# cluster_days_employed
def cluster_days_employed(x):
    days = x['DAYS_EMPLOYED']
    if days > 0:
        return 'not available'
    else:
        days = abs(days)
        if days < 30:
            return 'less 1 month'
        elif days < 180:
            return 'less 6 months'
        elif days < 365:
            return 'less 1 year'
        elif days < 1095:
            return 'less 3 years'
        elif days < 1825:
            return 'less 5 years'
        elif days < 3600:
            return 'less 10 years'
        elif days < 7200:
            return 'less 20 years'
        elif days >= 7200:
            return 'more 20 years'
        else:
            return 'not available'

```



```

df['cluster_days_employed'] = df.apply(cluster_days_employed, axis=1)

# custom_ext_source_3
def cluster_ext_source(x):
    if str(x) == 'nan':
        return 'not available'
    else:
        if x < 0.1:
            return 'less 0.1'
        elif x < 0.2:
            return 'less 0.2'
        elif x < 0.3:
            return 'less 0.3'
        elif x < 0.4:
            return 'less 0.4'
        elif x < 0.5:
            return 'less 0.5'
        elif x < 0.6:
            return 'less 0.6'
        elif x < 0.7:
            return 'less 0.7'
        elif x < 0.8:
            return 'less 0.8'
        elif x < 0.9:
            return 'less 0.9'
        elif x <= 1:
            return 'less 1'

df['cluster_ext_source_1'] = df['EXT_SOURCE_1'].apply(lambda x: cluster_ext_source(x))
df['cluster_ext_source_2'] = df['EXT_SOURCE_2'].apply(lambda x: cluster_ext_source(x))
df['cluster_ext_source_3'] = df['EXT_SOURCE_3'].apply(lambda x: cluster_ext_source(x))

# house_variables_sum
house_data = ['APARTMENTS_AVG', 'APARTMENTS_MEDI', 'APARTMENTS_MODE', 'BASEMENTAREA_AVG',
              'BASEMENTAREA_MEDI', 'BASEMENTAREA_MODE', 'COMMONAREA_AVG', 'COMMONAREA_MEDI',
              'COMMONAREA_MODE', 'ELEVATORS_AVG', 'ELEVATORS_MEDI', 'ELEVATORS_MODE', 'EMERGENCYSTATE_MODE',
              'ENTRANCES_AVG', 'ENTRANCES_MEDI', 'ENTRANCES_MODE', 'FLOORSMAX_AVG', 'FLOORSMAX_MEDI',
              'FLOORSMAX_MODE', 'FLOORSMIN_AVG', 'FLOORSMIN_MEDI', 'FLOORSMIN_MODE', 'FONDKAPREMONT_MODE',
              'HOUSETYPE_MODE', 'LANDAREA_AVG', 'LANDAREA_MEDI', 'LANDAREA_MODE', 'LIVINGAPARTMENTS_AVG',
              'LIVINGAPARTMENTS_MEDI', 'LIVINGAPARTMENTS_MODE', 'LIVINGAREA_AVG', 'LIVINGAREA_MEDI', 'LIVINGAREA_MODE',
              'NONLIVINGAPARTMENTS_AVG', 'NONLIVINGAPARTMENTS_MEDI', 'NONLIVINGAPARTMENTS_MODE', 'NONLIVINGAREA_AVG',
              'NONLIVINGAREA_MEDI', 'NONLIVINGAREA_MODE', 'TOTALAREA_MODE', 'WALLSMATERIAL_MODE',
              'YEARS_BEGINEXPLUATATION_AVG', 'YEARS_BEGINEXPLUATATION_MEDI', 'YEARS_BEGINEXPLUATATION_MODE',
              'YEARS_BUILD_AVG', 'YEARS_BUILD_MEDI', 'YEARS_BUILD_MODE']
df['house_variables_sum'] = df[house_data].sum(axis=1)
df['house_variables_sum_isnull'] = np.where(df['house_variables_sum'].isnull(), '1', '0')
df['house_variables_sum'].fillna(value=df['house_variables_sum'].median(), inplace=True)

num_columns = [
    'AMT_ANNUITY', 'AMT_CREDIT', 'AMT_GOODS_PRICE', 'AMT_INCOME_TOTAL',
    'REGION_POPULATION_RELATIVE', 'DAYS_BIRTH', 'DAYS_ID_PUBLISH', 'DAYS_REGISTRATION',
    'CNT_CHILDREN', 'CNT_FAM_MEMBERS', 'DAYS_EMPLOYED', 'DAYS_LAST_PHONE_CHANGE',
    'EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'total_amt_req_credit_bureau',
    'house_variables_sum']
cat_columns = [
    'CODE_GENDER', 'CSV_SOURCE', 'FLAG_OWN_CAR', 'NAME_EDUCATION_TYPE', 'FLAG_OWN_REALTY', 'OCCUPATION_TYPE', 'ORGANIZATION_TYPE',
    'NAME_CONTRACT_TYPE', 'NAME_FAMILY_STATUS', 'NAME_HOUSING_TYPE', 'NAME_INCOME_TYPE', 'NAME_TYPE_SUITE',
    'has_only_approved', 'has_been_rejected', 'has_job', 'has_children', 'cluster_days_employed',
    'cluster_ext_source_1', 'cluster_ext_source_2', 'cluster_ext_source_3',
    'total_amt_req_credit_bureau_isnull', 'house_variables_sum_isnull']

target_column = ['TARGET']
df = df[num_columns + cat_columns + target_column]

# Impute missing values
for numerical_column in num_columns:
    if df[numerical_column].isnull().values.any():
        df[numerical_column + '_isnull'] = np.where(df[numerical_column].isnull(), '1', '0')
        df[numerical_column].fillna(value=df[numerical_column].median(), inplace=True)

for categorical_column in cat_columns:
    df[categorical_column].fillna('NULL', inplace=True)

# Standard
minmax_scaler = preprocessing.MinMaxScaler()
df[num_columns] = pd.DataFrame(minmax_scaler.fit_transform(df[num_columns]))

```



```

#label_encoding
cat_columns.remove('CSV_SOURCE')

for column in cat_columns:
    df[column] = LabelEncoder().fit_transform(df[column].astype(str))
    df[column] = df[column].astype('category')

# split dataset
train_df = df[df['CSV_SOURCE'] == 'application_train.csv']
train_output_df = pd.DataFrame(train_df['TARGET'], columns=['TARGET'])

test_df = df[df['CSV_SOURCE'] == 'application_test.csv']

# remove columns
train_df.drop(columns=['CSV_SOURCE', 'TARGET'], axis=0, inplace=True)
test_df.drop(columns=['CSV_SOURCE', 'TARGET'], axis=0, inplace=True)

#validation
x_train, x_valid, y_train, y_valid = train_test_split(train_df, train_output_df, test_size=test_size, random_state=42)

def create_tensors(input_df):
    stack = []
    for column in input_df.columns:
        if input_df.dtypes[column] == np.int64 or input_df.dtypes[column] == np.float64:
            stack.append(input_df[column].astype(np.float64))
        else:
            stack.append(input_df[column].cat.codes.values)
    return torch.tensor(np.stack(stack, 1), dtype=torch.float)

tensor_x_train_cat = create_tensors(x_train[cat_columns]).float().to(device)
tensor_x_train_num = create_tensors(x_train[num_columns]).float().to(device)
tensor_y_train = torch.tensor(y_train.values).flatten().float().to(device)

tensor_x_valid_cat = create_tensors(x_valid[cat_columns]).float().to(device)
tensor_x_valid_num = create_tensors(x_valid[num_columns]).float().to(device)
tensor_y_valid = torch.tensor(y_valid.values).flatten().float().to(device)

tensor_x_test_cat = create_tensors(test_df[cat_columns]).float().to(device)
tensor_x_test_num = create_tensors(test_df[num_columns]).float().to(device)

# CREATE CATEGORICAL EMBEDDING SIZES
cat_columns_size = [len(df[column].cat.categories) for column in cat_columns]
categorical_embedding_sizes = [(col_size, min(50, (col_size + 1) // 2)) for col_size in cat_columns_size]

# DEFINE NEURAL NETWORK MODEL
class Model(nn.Module):
    def __init__(self, embedding_size, input_size, num_numerical_cols, layers, ps):
        super().__init__()

        self.all_embeddings = nn.ModuleList([nn.Embedding(ni, nf) for ni, nf in embedding_size])
        self.emb_drop = nn.Dropout(hp_emb_drop)

        self.bn_cont = nn.BatchNorm1d(num_numerical_cols)

        layer = []
        for i, elem in enumerate(layers):
            layer.append(nn.Linear(input_size, elem))
            layer.append(nn.ReLU(inplace=True))
            layer.append(nn.BatchNorm1d(layers[i]))
            layer.append(nn.Dropout(ps[i]))
            input_size = elem
        layer.append(nn.Linear(layers[-1], 1))

        self.layers = nn.Sequential(*layer)

    def forward(self, x_c, x_n):

        embeddings = [e(x_c[:,i]).long() for i, e in enumerate(self.all_embeddings)]

        x = torch.cat(embeddings, 1)
        x = self.emb_drop(x)

        x_n = self.bn_cont(x_n)

        x = torch.cat([x, x_n], 1)

```



```

    x = self.layers(x)

    return x

# INSTANCIATE MODEL

num_numerical_cols = tensor_x_train_num.shape[1]

num_categorical_cols = sum((nf for ni, nf in categorical_embedding_sizes))
initial_input_size = num_categorical_cols + num_numerical_cols

model = Model(categorical_embedding_sizes, initial_input_size, num_numerical_cols, layers=nnlayers, ps=hp_ps)
sigmoid = nn.Sigmoid()
loss_function = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=learningrate)
model.to(device)

# TRAIN NEURAL NETWORK MODEL
print("TRAINING MODEL...")
train_tensor_dataset = TensorDataset(tensor_x_train_cat, tensor_x_train_num, tensor_y_train)
train_loader = DataLoader(dataset=train_tensor_dataset, batch_size=batch_size, shuffle=True)

model.train()

tot_y_train_in = []
tot_y_train_out = []

for epoch in range(epochs):
    train_losses = []
    for x_cat, x_num, y in train_loader:
        y_train = model(x_cat, x_num)
        single_loss = loss_function(sigmoid(y_train.squeeze()), y)
        single_loss.backward()
        optimizer.step()

        train_losses.append(single_loss.item())
        tot_y_train_in.append(y)
        tot_y_train_out.append(y_train)
    epoch_loss = 1.0 * sum(train_losses) / len(train_losses)
    epoch_auc = roc_auc_score(torch.cat(tot_y_train_in).cpu().numpy(), torch.cat(tot_y_train_out).cpu().detach().numpy())
    tot_y_train_in = []
    tot_y_train_out = []
    print("\tepo: " + str(epoch) + "\tloss: " + str(epoch_loss) + "\tauc: " + str(epoch_auc))

# VALIDATE NEURAL NETWORK MODEL
print("VALIDATING MODEL...")
validation_tensor_dataset = TensorDataset(tensor_x_valid_cat, tensor_x_valid_num, tensor_y_valid)
validation_loader = DataLoader(dataset=validation_tensor_dataset, batch_size=batch_size, shuffle=True)

valid_losses = []

model.eval()

tot_y_valid_in = []
tot_y_valid_out = []

with torch.no_grad():
    for x_cat, x_num, y in validation_loader:
        y_valid = model(x_cat, x_num)
        validation_loss = loss_function(sigmoid(y_valid.squeeze()), y)
        valid_losses.append(validation_loss.item())

        tot_y_valid_in.append(y_valid)
        tot_y_valid_out.append(y)

valid_loss = round(1.0 * sum(valid_losses) / len(valid_losses), 5)
print("\tloss: " + str(valid_loss))
valid_auc = roc_auc_score(torch.cat(tot_y_valid_out).cpu(), torch.cat(tot_y_valid_in).cpu())
print("\tauc: " + str(valid_auc))

# MAKE PREDICTIONS
print("MAKING PREDICTIONS...")
with torch.no_grad():
    y_test = model(tensor_x_test_cat, tensor_x_test_num)

# GENERATE SUBMISSION.csv
print("GENERATING SUBMISSIONS...")

```

```

nn_prediction_df = pd.DataFrame(y_test.cpu().detach().numpy()).astype("float")
x_scaled = minmax_scaler.fit_transform(nn_prediction_df)
nn_prediction_df = pd.DataFrame(x_scaled)
nn_prediction_df = pd.concat([nn_prediction_df, application_test_df['SK_ID_CURR']], axis=1)
nn_prediction_df.columns = ['TEMP_TARGET', 'SK_ID_CURR']
nn_prediction_df['TARGET'] = nn_prediction_df['TEMP_TARGET']
nn_prediction_df = nn_prediction_df[['SK_ID_CURR', 'TARGET']]
nn_prediction_df.to_csv('submission1.csv', index=False)

print("EXECUTION COMPLETED.")

```

GPU - Tesla T4

```

<ipython-input-17-87fe9576a828>:121: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only'
df['house_variables_sum'] = df[house_data].sum(axis=1)

```


TRAINING MODEL...

epoch: 0	loss: 0.7139497920967668	auc: 0.6309687141142222
epoch: 1	loss: 0.667470389984364	auc: 0.6886699362857125
epoch: 2	loss: 0.6343034151626656	auc: 0.7217084690823279
epoch: 3	loss: 0.574318533523495	auc: 0.7253117017023164
epoch: 4	loss: 0.5002443059523177	auc: 0.7288723717833956
epoch: 5	loss: 0.4283307173782269	auc: 0.7368439610165431
epoch: 6	loss: 0.36568192538886446	auc: 0.7429465721912977
epoch: 7	loss: 0.31617653191632195	auc: 0.7477063433327675
epoch: 8	loss: 0.28012465405603687	auc: 0.7544540997999702
epoch: 9	loss: 0.25712734706922685	auc: 0.7631410760378139
epoch: 10	loss: 0.2450951326769257	auc: 0.7690683618517538
epoch: 11	loss: 0.24061422689450887	auc: 0.7762495197384395

VALIDATING MODEL...

loss: 0.25533
auc: 0.738632800087059

MAKING PREDICTIONS...



GENERATING SUBMISSIONS...

EXECUTION COMPLETED.

model

```

Model(
  (all_embeddings): ModuleList(
    (0): Embedding(3, 2)
    (1): Embedding(2, 1)
    (2): Embedding(5, 3)
    (3): Embedding(2, 1)
    (4): Embedding(19, 10)
    (5): Embedding(58, 29)
    (6): Embedding(2, 1)
    (7-8): 2 x Embedding(6, 3)
    (9-10): 2 x Embedding(8, 4)
    (11-12): 2 x Embedding(3, 2)
    (13-14): 2 x Embedding(2, 1)
    (15): Embedding(9, 5)
    (16): Embedding(11, 6)
    (17-18): 2 x Embedding(10, 5)
    (19): Embedding(2, 1)
  )
)

```



```

        (20): Embedding(1, 1)
    )
    (emb_drop): Dropout(p=0.04, inplace=False)
    (bn_cont): BatchNorm1d(17, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (layers): Sequential(
      (0): Linear(in_features=107, out_features=800, bias=True)
      (1): ReLU(inplace=True)
      (2): BatchNorm1d(800, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (3): Dropout(p=0.001, inplace=False)
      (4): Linear(in_features=800, out_features=350, bias=True)
      (5): ReLU(inplace=True)
      (6): BatchNorm1d(350, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (7): Dropout(p=0.01, inplace=False)
      (8): Linear(in_features=350, out_features=1, bias=True)
    )
  )

layer_sizes = []
for module in model.layers:
    if isinstance(module, nn.Linear):
        layer_sizes.append(module.out_features)

# Create the architecture string
architecture_string = '-'.join(map(str, layer_sizes))
print(f'architecture_string', architecture_string)

architecture_string 800-350-1

##score
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score

y_true = torch.cat(tot_y_valid_out).cpu().numpy()
y_pred_probs = torch.cat(tot_y_valid_in).cpu().numpy()

y_pred = (y_pred_probs > 0.5).astype(int)

accuracy = accuracy_score(y_true, y_pred)
recall = recall_score(y_true, y_pred)
precision_score = precision_score(y_true, y_pred)
f1_score = f1_score(y_true, y_pred)


print(f'Accuracy: {accuracy:.4f}')
print(f'f1_score: {f1_score:.4f}')
print(f'precision: {precision_score:.4f}')
print(f'Recall: {recall:.4f}')

Accuracy: 0.9194
f1_score: 0.0263
precision: 0.5076
Recall: 0.0135

# Calculate AUC-ROC and plot the curve
fpr, tpr, _ = roc_curve(torch.cat(tot_y_valid_out).cpu(), torch.cat(tot_y_valid_in).cpu())
roc_auc = auc(fpr, tpr)

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange', lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```

**Discussion:**

Architecture string of Neural Networks - 800-350-1, 246-32-16-8-2

Architecture string used - 800-350-1 Multiple architectures were used to build a optimized Neural Networks and have choose neural network with accuracy was 91% with auc curve - 0.74, f1_score: 0.0263

0.0 0.2 0.4 0.6 0.8 1.0

Key Experiments✓ **Logistic Regression**

```
#Importing Libraries
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
```

```
#Loading dataset
df_clean = pd.read_csv('clean_final.csv')
```

```
df1 = df_clean.copy()
```